

Nordea

Migrating accounting flows - from legacy platforms to Big Data

A short introduction

Goran Amcoff, BI process specialist



Nordea Bank – at glance

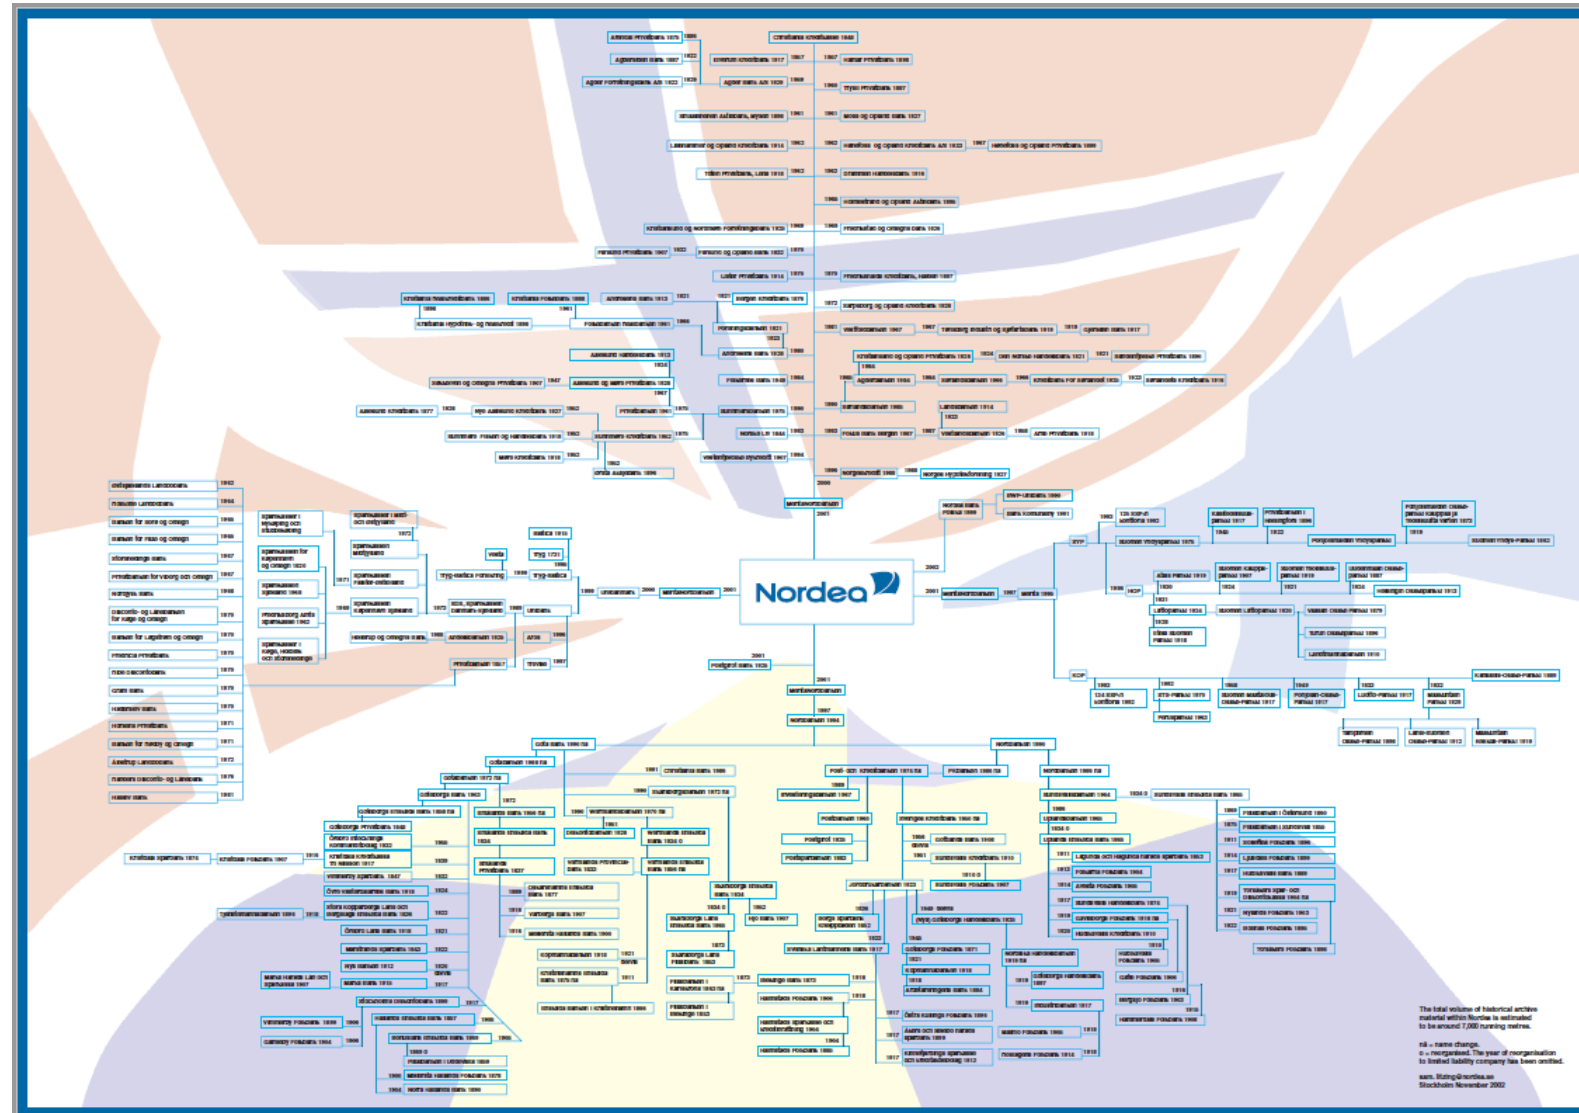
- Nordic financial services group operating from Helsinki
- >32000 employees
- >9 billion EUR revenue
- ~11 million private and corporate customers
- >600 billion \$ in total assets
- Major presence in Sweden, Norway, Denmark and Finland



Mergers & acquisitions present unique challenges for accounting flows

- Mergers & Acquisitions = inheriting a lot of technology from the companies you are merging with.
- It is not easy to merge mature organizations into common accounting system.
- Accounting must work from day 1, but many things are done differently.
- Pragmatic approach: countries/branches send accounting through pre-defined API. Integration is taken care by “abstraction layer” = Nordea GAI 1.0.

Nordea family tree



Current accounting solution – GAI 1.0

- Pragmatic solution for complicated landscape: General Accounting Interface - GAI 1.0
- Countries/Branches run different core systems, but are all delivering data in common format.
- Master Data is used for lookups and/or translation of dissimilar accounts
- Data ingest and light transformation done by TWS/Data Stage.
- Heavy logic/transformation is done in COBOL.

PROS and CONS of legacy GAI 1.0 solution

Mainframe is proven and robust technology

Good single thread performance.

Stable suppliers that have been around for a long time

Processes in place (debugging, testing, deployment etc.) have been rehearsed for years or even decades.

Pricy

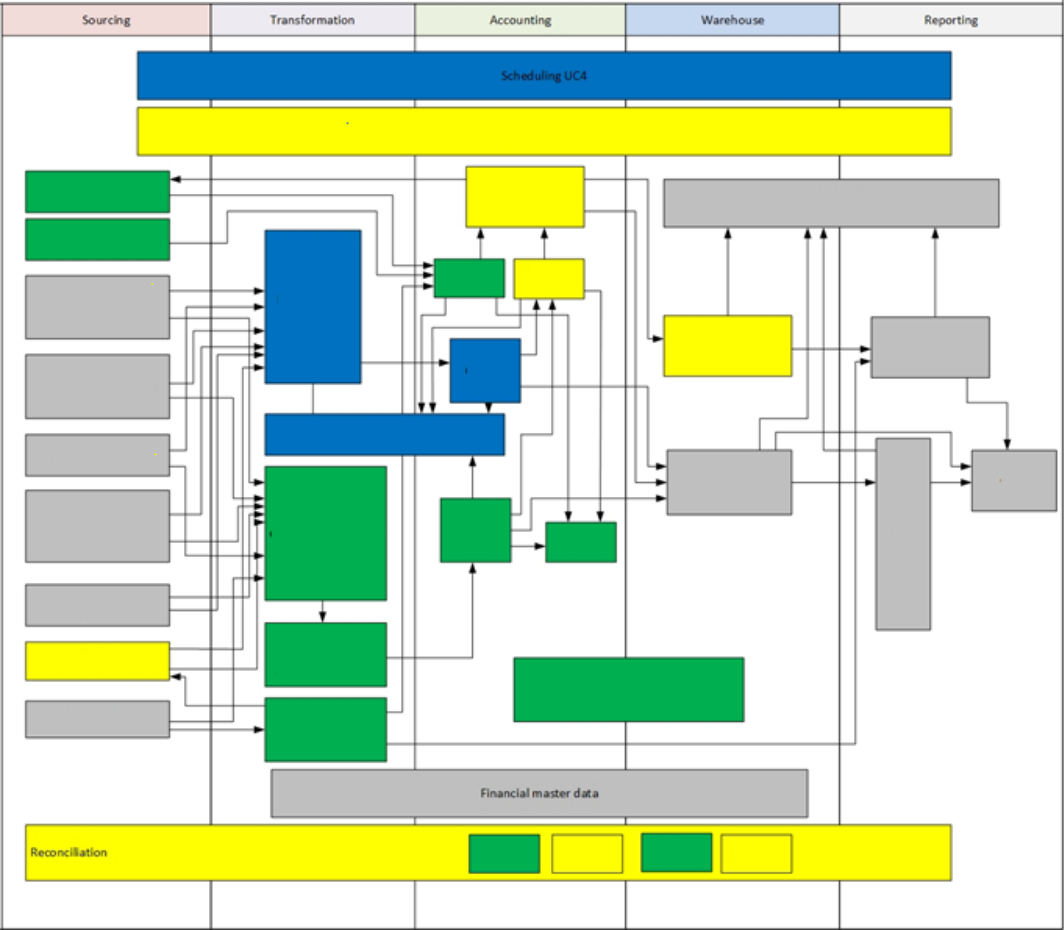
Slow change process

Securing resources with MF skills is getting harder

Horizontal scaling not easy

Most of tech stack is closed source

Accounting is complicated



What is driving migration from GAI1

- Number of transactions in core banking systems is continuously growing. We plan for >300Mtrans/day processed within hours
- Certain operations present unique challenges: Markets produce prodigious amount of transactions.
- Legal requirements: we are required to store all transactions for 10 years, including meta-data. Archive grows continuously and storage on MF is pricey.
- People who built GAI 1.0 and understand its intricacies are retiring.
- Largest performance driver: Treasury is pushing for larger implementation of **FTP** (Funds Transfer Pricing).
https://en.wikipedia.org/wiki/Funds_transfer_pricing

FTP

- A great tool for financial performance evaluation
- ...also an formidable IT issue
- FTP can introduce a large multiplier on amount of transactions. What was 10M transactions might quickly becomes 500M transactions.
- Eats lots of CPU, requires lookup tables and consumes storage for all meta data.

New accounting platform feature list

- Audit Trail
- Transaction processing

- Newer technology (Java? x86?) * *
- Horizontally scalable in with small granularity (adding cheap nodes vs. adding new racks/10 years contracts)
- Possible to migrate to “cloud” if/when needed * *
- Able to **cheaply and reliably** store and do light processing/retrieval of vast amounts of data/meta-data (Multi Petabyte)
- Easier to use modern UI libraries (React, node.js etc) * *
- Open Source building blocks * *

GAI2 architecture

- UC4 instead of Data Stage/TWS
- Java instead of COBOL for calculating engines
- WebLogic/React for UI layers (replacing legacy UI)
- Hadoop cluster for archival storage (replacing DB2/MF datasets)
- Hadoop cluster for landing areas (replacing SAN's and MF)
- Keep RDBMS at minimum
- REST-protocol for M2M communication (replacing scheduled FTP/flat files/ODBC etc.)
- Hadoop for Data Quality checks

Challenges/lessons re: GAI2

- You still need relational databases, Hadoop alone is not enough. (Joins are easier in SQL than Java/Map-Reduce).
- Developers should work near process specialists and SME's. Otherwise they might assume things. (...what needs to be logged and what can be discarded)
- If it works in demo it might not work with production volumes. (Java garbage collection, performance, DB logs etc).
- Performance profiling: system must be able to handle volumes within reasonable time window...
- Optimization skills are still very much needed, some things are not solvable by more hardware.

Showcase for Big Data tech stack: TRAT2 (TRansaction Archive Tool 2.0)

TRAT2 is tool for tracking and storing accounting data. It is meant to replace TRAT1.

Legacy: TRAT1 - uncompressed datasets on mainframe, Web-based UI. Data Stage/TWS used for ETL, flat files for M2M communication.

- Business search (search for any field) not possible.
- Search speed varies, usually minutes.
- Search for historical data very slow (>20 minutes)
- Not easily scalable performance wise. Scalable capacity-wise.
- Old data "nearline", not "online".

New: TRAT2- Cloudera Hadoop cluster, Apache Hive, Apache SOLR indexer, React/Node.js for UI. UC4 for ETL. REST services for Hive<->Web Services communication.

- Business search possible , logical and text operators (LIKE, <,> etc.) possible
- 50 - 100 times (!) compression ratio compared to TRAT1 due to data being highly formalized and repetitive.
- Horizontally scalable. Need more storage/CPU? Add more nodes...
- Schema on read: Transaction files are stored "as is".
- Huge performance increases when it comes to search times for indexed "hot data" (<3 sec).
- Relatively good search times for "cold data" (<5 mins).
- All data is "online".

Challenges re:TRAT2

- Which Hadoop release/version/package to use? Cloudera? Map-R?
- Hadoop concept differs from legacy hardware. Infrastructure people must understand it (lot's of small/simple nodes instead of expensive/powerful servers.)
- Making Hadoop resilient: node replication factor? Failover strategy? Offline backup?
- How do you display large volume of data through React-based UI without crashing the browser or running out of memory on server side?
- Indexing vs. “cold search”: how many generations of data are to be indexed (=requires extra storage/processing) vs. just stored and searchable via “table scan”
- Ingestion/migration of >1PB from legacy system (TRAT1):
 - Do you convert data to “intermediate” format?
 - How to transport 10^{15} bytes and make sure data is intact (checksums, data sampling)
 - Ingesting and indexing 10^{15} bytes...

Do not forget about people doing the work!

- People management is as important as technology.
- People working with “legacy” systems are just as important as those developing new system! **Who will tell you how it works?** Avoid “Us and Them”.
- Create cross-functional teams. Do not forget to involve business people. They might not know much about technology but know a lot about what system should do and what is important and what is not.
- Beware of vendor lock-in and outsourcing both architecture and execution. Key stakeholders/drivers should come from the organization.



Q/A