

Regularizing Black-box Models for Improved Interpretability

Gregory Plumb¹ Maruan Al-Shedivat¹ Ángel Alexander Cabrera¹ Adam Perer¹ Eric Xing¹
Ameet Talwalkar¹

Abstract

Most of the work on interpretable machine learning has focused on designing either inherently interpretable models, which typically trade-off accuracy for interpretability, or post-hoc explanation systems, which tend to lack guarantees about the quality of their explanations. We explore a hybridization of these approaches by directly regularizing a black-box model for interpretability at training time—a method we call ExpO. We find that post-hoc explanations of an ExpO-regularized model are consistently more stable and of higher fidelity, which we show theoretically and support empirically. Critically, we also find ExpO leads to explanations that are more actionable, significantly more useful, and more intuitive as supported by a user study.

1. Introduction

Complex learning-based systems are increasingly shaping our daily lives. To monitor and understand these systems, we require clear explanations of model behavior. While model interpretability has many definitions and is often largely application specific (Lipton, 2016), local explanations are a popular and powerful tool (Ribeiro et al., 2016). Recent work on interpretability in machine learning ranges from proposals of new models that are interpretable *by-design* (e.g., Wang & Rudin, 2015; Caruana et al., 2015) to model-agnostic *post-hoc* algorithms for explaining complex, black-box predictors such as ensembles and deep neural networks (e.g., Ribeiro et al., 2016; Lei et al., 2016; Lundberg & Lee, 2017; Selvaraju et al., 2017; Kim et al., 2018). Despite the variety of technical approaches, the underlying goal of all these works is to develop an interpretable predictive system that produces two outputs: a prediction and an explanation, which can be useful for reasoning about the system as well as for making downstream decisions.

Both interpretability by-design and post-hoc explanation

¹Carnegie Mellon University. Correspondence to: Gregory Plumb <gdplumb@andrew.cmu.edu>.




Figure 1: Neighborhood Fidelity of LIME-generated explanations (lower is better) vs. predictive error of several models trained on the UCI ‘housing’ regression dataset. The values in blue denote regularization weight of ExpO.

strategies have limitations. On the one hand, the by-design approaches are restricted to working with model families that provide inherent interpretability, potentially at the cost of accuracy. On the other hand, by performing two disjoint steps, there is no guarantee that post-hoc explainers applied to an arbitrary model will produce explanations of suitable quality. Moreover, recent approaches that claim to overcome this apparent trade-off between prediction accuracy and explanation quality are in fact by-design proposals that impose certain constraints on the underlying model families they consider (e.g., Al-Shedivat et al., 2017; Plumb et al., 2018; Alvarez-Melis & Jaakkola, 2018a).

In this work, we propose a strategy called *Explanation-based Optimization* (ExpO) that allows us to interpolate between the two paradigms by adding an *interpretability regularizer* to the loss function of an arbitrary predictive model. We illustrate how ExpO can influence the interpretability and accuracy of a model in Figure 1.

Illustration. Consider a situation where Bob’s house price is estimated by a machine learning system. Here, a good local explanation (*i.e.*, an explanation specific to Bob’s situation) should help Bob understand what factors influenced his house’s evaluation and what he can do to increase its predicted value. Unfortunately, a standard multi-layer perceptron model trained with SGD is often difficult to explain. To improve the local explanations of this model, we propose

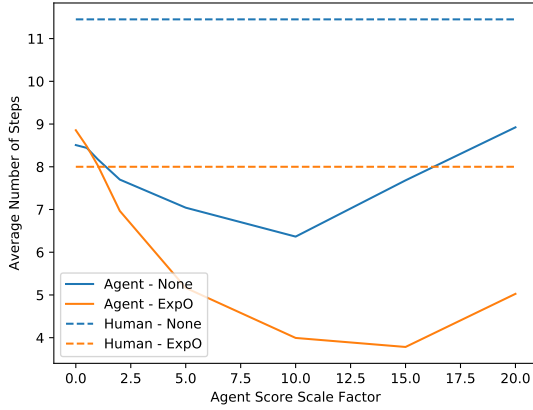


Figure 2: A comparison of the number of steps it takes either a human or an algorithmic agent to use a model’s explanations to complete the task described in Section 5. The x-axis is a measure of the agent’s randomness: 0 being totally random with the agent becoming greedier with larger values (See the Appendix A.1 for a detailed definition). Note that both the humans and the agent found it easier to complete this task for a model trained with ExpO. Interestingly, humans actually perform worse than random guessing for a normally trained model.

to train it with regularizers based on the *fidelity* (Ribeiro et al., 2016; Plumb et al., 2018) or *stability* (Alvarez-Melis & Jaakkola, 2018a) metrics that measure the quality of local explanations. By directly optimizing these metrics, we intend to bias the model towards being more amenable to local explanation. To evaluate the effectiveness of our approach, we run a user study and gauge the usefulness of the information provided by the explanations to aid human decisions similar to Bob’s (Figure 2).

The specific contributions of our work are as follows:

1. **Interpretability regularizers.** We introduce two regularizers—ExpO-Fidelity and ExpO-Stability—that are designed to improve model interpretability with respect to the fidelity and stability metrics, respectively.
2. **Generalizable explanation quality.** We show that the benefits of our regularization technique generalize to unseen points with respect to explanation quality metrics. Specifically, we bound on the gap between the fidelity of explanations on training and held out points, connecting it with the local variance of the learned model.
3. **Empirical results.** We evaluate models trained with and without our regularizers on a variety of regression and classification tasks.¹ Empirically, regularization does not affect (or slightly improves) the predictive performance across the nine datasets we consider, while showing significant improvement in terms of explanation quality

¹<https://github.com/GDPlumb/ExpO>

(at least 25% improvement in terms of fidelity).

4. **User study.** Although helpful proxies, the fidelity and stability metrics are only indirect measures of the usefulness of local explanations. To test the validity of ExpO in a real-world scenario, we ran a user study with 60 participants, asking them to change a model’s input in order to get the predicted value into a specific range using information from local explanations. We found that participants both preferred the explanations from the ExpO regularized model and were able to complete the task in less steps than with the unregularized model.

2. Background and Related Work

Consider a supervised learning problem, where our goal is to estimate a model, $f : \mathcal{X} \mapsto \mathcal{Y}$, $f \in \mathcal{F}$, that maps input feature vectors, $x \in \mathcal{X}$, to targets, $y \in \mathcal{Y}$, and is trained using data, $\{x_i, y_i\}_{i=1}^N$. If the class of functions used for modeling the data is complex, we can understand the behavior of f in some neighborhood, $N_x \in \mathcal{P}[\mathcal{X}]$ (where $\mathcal{P}[\mathcal{X}]$ is the space of probability distributions over \mathcal{X}), by generating a local *explanation*.

We denote algorithms that produce local explanations (*i.e.*, *explainers*) as $e : \mathcal{X} \times \mathcal{F} \mapsto \mathcal{E}$, where \mathcal{E} is the set of possible explanations. The choice of \mathcal{E} generally depends on whether or not \mathcal{X} consists of *semantic features*, and will be defined more precisely next.

2.1. Semantic Features

We call features *semantic* if people can reason about them and understand what it means when their values change (*e.g.*, a person’s income, the concentration of a chemical, etc.). Consequently, local explanations try to predict how the model’s output would change if the input was perturbed (Ribeiro et al., 2016; Plumb et al., 2018). Thus, we can define the output space of the explainer as $\mathcal{E}_s := \{g \in \mathcal{G} \mid g : \mathcal{X} \mapsto \mathcal{Y}\}$, where \mathcal{G} is a class of interpretable (typically linear) functions.

Fidelity metric. When the explainer’s output space is \mathcal{E}_s , the explanation is defined as a function $g : \mathcal{X} \mapsto \mathcal{Y}$, and it is natural to evaluate how accurately g models f in a neighborhood N_x (Ribeiro et al., 2016; Plumb et al., 2018):

$$F(f, g, N_x) := \mathbb{E}_{x' \sim N_x} [(g(x') - f(x'))^2], \quad (1)$$

which we refer to as the *neighborhood-fidelity* (NF) metric. This metric is sometimes evaluated with N_x as a point mass on x and we call this version the *point-fidelity* (PF) metric. While Plumb et al. (2018) argued that point-fidelity can be misleading because it does not measure generalization of $e(x, f)$ across N_x , it has been used for evaluation in the prior work (Ribeiro et al., 2016; 2018) and we report it in

our experiments along with the neighborhood-fidelity for completeness.

Black-box explanation systems. Various explainers have been proposed to generate local explanations of the form $g : \mathcal{X} \mapsto \mathcal{Y}$, typically assuming that g is linear. In particular, LIME (Ribeiro et al., 2016), one of the most popular black-box explanation systems, solves the following optimization problem:

$$e(x, f) := \arg \min_{g \in \mathcal{E}_s} F(f, g, N_x) + \Omega(g), \quad (2)$$

where $\Omega(e)$ stands for an additive regularizer that encourages certain desirable properties of the explanations (e.g., sparsity). LIME’s objective function is closely related to the fidelity metric and subsequently to our proposed ExpO-Fidelity regularizer. Consequently, we expect our regularizer to improve the quality of LIME-generated explanations. Our experimental results in Section 4.1 corroborate this hypothesis.

Along with LIME, we consider another black-box explanation tool, called MAPLE (Plumb et al., 2018). It differs substantially from LIME in that its *neighborhood function is learned from the data* rather than specified as a parameter. In our experiments, we evaluate the quality of MAPLE-generated local explanations for models regularized via ExpO-Fidelity, but do not use MAPLE’s learned neighborhood function to define ExpO-Fidelity. We view this as a good test case to see how optimizing the fidelity metric for one neighborhood generalizes to another one (see Section 3 for a more detailed discussion of this point). In Section 4.1, we see that regularizing for LIME neighborhoods improves MAPLE’s explanation quality as well.

2.2. Non-Semantic Features

Non-semantic features lack an inherent interpretation, with images being a canonical example. When \mathcal{X} consists of non-semantic inputs, we cannot assign meaning to the difference between x and x' , hence it does not make sense to explain the difference between the predictions $f(x)$ and $f(x')$. As a result, fidelity is not an appropriate explanation metric. Instead, in this context, local explanations try to identify which parts of the input are particularly influential on a prediction (Lundberg & Lee, 2017; Sundararajan et al., 2017). Consequently, we consider explanations of the form $\mathcal{E}_{ns} := \mathbb{R}^d$, where d is the number of features in \mathcal{X} .

Stability metric and saliency maps. When the explainer’s output space is \mathcal{E}_{ns} , the explanation is a vector in \mathbb{R}^d , and cannot be directly compared to the underlying model itself, as in the case of the fidelity metric. Instead, the focus in this setting is on the degree to which the explanation changes between points in a local neighborhood,

which we measure using the *stability metric* (Alvarez-Melis & Jaakkola, 2018a):

$$\mathcal{S}(f, e, N_x) := \mathbb{E}_{x' \sim N_x} [\|e(x, f) - e(x', f)\|_2^2] \quad (3)$$

Various explainers (Sundararajan et al., 2017; Zeiler & Fergus, 2014; Shrikumar et al., 2016; Smilkov et al., 2017; Adebayo et al., 2018) have been proposed to generate local explanations in \mathcal{E}_{ns} , with *saliency maps* (Simonyan et al., 2013) being the approach that we consider in this work. Saliency maps assign importance weights to image pixels based on the magnitude of the gradient of the predicted class with respect to the corresponding pixels.

Recent work on model interpretability emphasizes that more stable explanations tend to be more trustworthy (Alvarez-Melis & Jaakkola, 2018a; Ghorbani et al., 2017; Alvarez-Melis & Jaakkola, 2018b). Note that the stability metric can also be considered in the context of semantic features in addition to the fidelity metric.

2.3. Related Methods

A few recently proposed approaches to model interpretability are closely related to our work. First, the work most closely related to ours is (Lee et al., 2018; 2019), which explored regularizing black-box models to be more locally interpretable for time-series and graph data. Second, self-explaining neural networks (SENN) (Alvarez-Melis & Jaakkola, 2018a) (a variation of contextual explanation networks (Al-Shedivat et al., 2017)) is an interpretable by-design approach that additionally (indirectly) optimizes their models to produce stable explanations. Third, “Right For The Right Reasons” (RTFR) (Ross et al., 2017) selectively penalizes gradients of the output with respect to certain input features at some points to discourage their use by the model.

Because ExpO and the method from (Lee et al., 2018; 2019) are designed to work with different types of data, there are a few key differences between them. Primarily, the regularizers used in (Lee et al., 2018; 2019) are non-differentiable while those in ExpO are differentiable. Next, the time-series definition of the neighborhood in (Lee et al., 2018; 2019) looks at a different type of locally explainability that focuses on making it easier to understand what the model will predict for the next point in the series rather than how its prediction would change if the time series itself was changed, which is the focus of this work. Finally, the graph definition of the neighborhood in (Lee et al., 2019) requires domain knowledge to define while ExpO’s regularization does not.

From a technical standpoint, SENN and RTFR both assume that the local explanation is close to the first order Taylor approximation of the model at that point. In Section 3.1, we demonstrate how Taylor approximations are often quite different from and more difficult to use than the

neighborhood-based local explanations that we use in ExpO. Further, SENN’s regularizer requires the neural network to have a very particular structure and, therefore, unlike ExpO, cannot be applied to an arbitrary model. While RTFR’s regularization can be used with arbitrary models, it is not directly related to a measure of explanation quality and is defined using specific domain knowledge; on the other hand, ExpO aims to directly improve quality of explanations with respect to a specific metric and does not require domain knowledge.

In the Appendix A.2, we compare ExpO to simple l_1 and l_2 regularization since other baselines are either model specific or require domain knowledge. We found that these strategies do not significantly impact model interpretability.

3. Explanation Optimization

Running black-box explainers on arbitrary models does not guarantee the quality of the resulting explanations. To address this, we define regularizers that can be added to the loss function and used to train an arbitrary model f . Specifically, we want to solve the following optimization problem:

$$\hat{f} := \arg \min_{f \in \mathcal{F}} \frac{1}{N} \sum_{i=1}^N (\mathcal{L}(f, x_i, y_i) + \gamma \mathcal{R}(f, N_{x_i}^{\text{reg}})) \quad (4)$$

where $\mathcal{L}(f, x_i, y_i)$ is a standard predictive loss (e.g., squared error for regression or cross-entropy for classification), $\mathcal{R}(f, N_{x_i}^{\text{reg}})$ is a regularizer that encourages explainability of f in the neighborhood of x_i , and $\gamma > 0$ controls the regularization strength. Because our regularizers are differentiable, we can solve Equation 4 using any standard gradient-based algorithm; in our case, SGD with Adam.

We define $\mathcal{R}(f, N_x^{\text{reg}})$ based on either the neighborhood-fidelity, Eq. (1), or the neighborhood-stability, Eq. (3). In order to compute these metrics exactly, we would need to run an explainer algorithm, e ; this may be non-differentiable or too computationally expensive to use as a regularizer. Thus, for ExpO-Fidelity, we approximate e using a local linear model fit on points sampled from N_x^{reg} (Algorithm 2). For ExpO-Stability, we simply require that the model’s output not change too much across N_x^{reg} (Algorithm 3).²

Choosing the neighborhood. Defining a good regularization neighborhood, requires considering the following. On the one hand, we would like N_x^{reg} to be similar to N_x , as used in Eq. 1 or Eq. 3, so that the neighborhoods used for regularization and for evaluation match. On the other hand, we also would like N_x^{reg} to be consistent with the local neighborhood defined by e internally, which may differ from N_x . For LIME, this is not a problem since the internal

definition of the local neighborhood is a hyperparameter that we can set. However for MAPLE, the local neighborhood is learned from the data, and hence the regularization and explanation neighborhoods may differ. Ultimately, we left resolving this tension to future work.

Computational cost. Algorithm 2 could be prohibitively expensive since the number of samples, m , from N_x^{reg} , has to be proportional to the dimension of x , d , resulting in $O(d^3)$ evaluations of f to compute the regularizer at x . So we also test a randomized version of Algorithm 2, ExpO-1D-Fidelity, that randomly selects one dimension of x to perturb according to N_x^{reg} and penalizes the error of a local linear model along that dimension. This breaks the dependence of the computational cost of the regularizer on d and allows us to compute each gradient step with a $O(1)$ increase in the number of evaluations of f .

3.1. Understanding the Properties of ExpO

The goal of this section is to compare the behavior of local linear explanations and our regularizers to some existing theoretical function approximations and measures of variance to help develop an intuitive understanding of ExpO. First, we compare neighborhood-based local linear explanations to first order Taylor approximations to show that they can have fundamentally very different behaviors. Second, we compare ExpO-Fidelity to the Lipschitz Constant (LC) and Total Variation (TV) of the learned function.

Local explanations vs. Taylor approximations. A natural question to ask is, *Why should we sample from N_x in order to locally approximate f when there are easier and theoretically motivated approximations?* One possible way to do this is via the Taylor approximation (Alvarez-Melis & Jaakkola, 2018a). The downside of a Taylor approximation-based approach is that such an approximation cannot readily be adjusted to different neighborhood scales and its fidelity and stability strictly depend on the learned function. This can be seen in Figure 3 where the Taylor approximations at two nearby points are both radically different and not faithful to the model outside of an small neighborhood.

Fidelity regularization and the model’s LC or TV. From a theoretical perspective, our regularizer is similar to controlling the Lipschitz Constant or Total Variation of f across N_x after removing the part of f explained by $e(x, f)$. From an interpretability perspective, there is nothing inherently wrong with having a large LC or TV, which is demonstrated in Figure 4. However, once we take into account what can be explained by $e(x, f)$, then upper bounding any one of ExpO-Fidelity, the LC, or the TV will upper bound the others.

²A similar procedure was explored previously in (Zheng et al., 2016) for adversarial robustness.

Algorithm 1 Learning with ExpO

input \mathcal{D} – dataset, γ – regularization coefficient,
 N_x^{reg} – neighborhood sampling function,
 T – number of optimization epochs.

1: Initialize model f with parameters θ_0 .

2: **for** t in $0 \dots T$ **do**

3: **for** each of mini-batch \mathcal{B} in \mathcal{D} **do**

4: **for** each point x_i in \mathcal{B} **do**

5: Regularizer: $r(\theta, x_i) = \mathcal{R}(f_\theta, N_{x_i}^{\text{reg}})$
 (compute using Algorithm 2 or 3).

6: **end for**

7: Construct loss:

$$L(\theta) := \frac{1}{|\mathcal{B}|} \sum_{x_i \in \mathcal{B}} \mathcal{L}(\theta, x_i) + \gamma r(\theta, x_i).$$

8: Update the model:

$$\theta_{t+1} \leftarrow \text{update}(\theta_t, \nabla L(\theta_t))$$

9: **end for**

10: **end for**

output Learned model, f_{θ_T} .

3.2. Generalization of Local Linear Explanations

To conclude our analysis, we study the quality of local linear explanations in terms of generalization. Note that ExpO regularization encourages learning models that are explainable in the neighborhoods of each *training point*. However, how would this property generalize to unseen points? We answer this question by providing a generalization bound in terms of neighborhood-fidelity metric for local linear explanations (see Appendix A.3 for derivations).

Proposition 1 *Let the neighborhood sampling function N_x be characterized by some parameter σ (e.g., the effective radius of a neighborhood) and the variance of the trained model $f(x)$ across all such neighborhoods be bounded by some constant $C(\sigma) > 0$. Then, the following bound holds with at least $1 - \delta$ probability:*

$$\mathbb{E} [r(f, x)] \leq \frac{1}{n} \sum_{i=1}^n r(f, x_i) + \sqrt{\frac{C^2(\sigma) \log \frac{1}{\delta}}{2n}} \quad (5)$$

Remark 2 *The obtained bound tells us that models with smaller local variances across the neighborhoods are likely to have explanations of higher fidelity on the held out points. This further motivates the approximation we used in Algorithm 3.*

Algorithm 2 Neighborhood-fidelity regularizer

input $f_\theta, x, N_x^{\text{reg}}, m$

1: Sample points: $x'_1, \dots, x'_m \sim N_x^{\text{reg}}$

2: Compute predictions:

$$\hat{y}_j(\theta) = f_\theta(x'_j) \text{ for } j = 1, \dots, m$$

3: Produce a local linear explanation:

$$\beta_x(\theta) = \arg \min_{\beta} \sum_{j=1}^m (\hat{y}_j(\theta) - \beta^\top x'_j)^2$$

output $\frac{1}{m} \sum_{j=1}^m (\hat{y}_j(\theta) - \beta_x(\theta)^\top x'_j)^2$

Algorithm 3 Neighborhood-stability regularizer

input $f_\theta, x, N_x^{\text{reg}}, m$

1: Sample points: $x'_1, \dots, x'_m \sim N_x^{\text{reg}}$

2: Compute predictions:

$$\hat{y}_j(\theta) = f_\theta(x'_j), \text{ for } j = 1, \dots, m$$

output $\frac{1}{m} \sum_{j=1}^m (\hat{y}_j(\theta) - f(x))^2$

4. Experimental Results

In our first set of experiments, we demonstrate the effectiveness of ExpO-Fidelity and ExpO-1D-Fidelity on datasets with semantic features using seven regression problems from the UCI collection (Dheeru & Karra Taniskidou, 2017) as well as an in-hospital mortality classification problem. Dataset statistics are in Table 1. Our second experiment demonstrates the effectiveness of ExpO-Stability for creating saliency maps (Simonyan et al., 2013) on MNIST (LeCun, 1998). We found that a model trained with our regularizers is more interpretable than a model trained without them because black-box explainers produce quantitatively better explanations for them; further, they are often more accurate. Finally, we demonstrate qualitatively that the explanations for the regularized model tend to be simpler.

4.1. Neighborhood-Fidelity Regularization

We compare models trained with our regularizers to models trained without them. We report accuracy and three interpretability metrics: Point-Fidelity (PF), Neighborhood-Fidelity (NF), and Stability (S). The interpretability metrics are evaluated for two black-box explanation systems: LIME and MAPLE. Consequently, the “MAPLE-PF” label corresponds to the Point-Fidelity Metric for explanations produced by MAPLE.

Experimental setup. All of the inputs to the model were standardized to have mean zero and variance one (including the response variable for regression problems). The network




Figure 3: A function (blue), its first order Taylor approximations at $x = 0.4$ (green) and $x = 0.5$ (red), and a local explanation of the function (orange) computed with $x = 0.5$ and $N_x = [0, 1]$.




Figure 4: Top: Two functions (blue) and their local linear explanations (orange). The local explanations were computed with $x = 0.5$ and $N_x = [0, 1]$. **Bottom:** The unexplained portion of the function (residuals).

Table 1: Statistics of the datasets.

Dataset	# samples	# dims
autmpgs	392	7
communities	1993	102
day	731	14
housing	506	11
music	1059	69
winequality-red	1599	11
MSD	515345	90
SUPPORT2	9104	51
MNIST	60000	784

Table 2: An example of LIME’s explanation coefficients for unregularized and ExpO-regularized models.

Feature	x	Unreg.	ExpO
CRIM	2.5	-0.05	-0.03
INDUS	1.0	0.1	-0.01
NOX	0.9	-0.23	-0.18
RM	1.4	0.22	0.2
AGE	1.0	-0.08	0.02
DIS	-1.2	-0.38	-0.15
RAD	1.6	0.24	0.17
TAX	1.5	-0.27	-0.11
PTRATIO	0.8	-0.11	-0.14
B	0.4	0.12	-0.01
LSTAT	0.1	-0.34	-0.53




Figure 5: Original images (left) and saliency maps of an unregularized (middle) and regularized (right) models.

architectures and hyper-parameters were chosen using a simple grid search. For the final results, we set N_x to be $\mathcal{N}(x, \sigma)$ with $\sigma = 0.1$ and N_x^{reg} to be $\mathcal{N}(x, \sigma)$ with $\sigma = 0.5$. In the Appendix A.4, we discuss how these values were chosen.

UCI regression experiments. The effects of ExpO-Fidelity and ExpO-1D-Fidelity on model accuracy and interpretability are in Table 3. ExpO-Fidelity frequently improved the interpretability metrics by over 50%, with the smallest improvements being around 25%. Further, it lowered the prediction error on the ‘communities’, ‘day’, and ‘MSD’ datasets, which lets us conclude that it has a small positive effect on accuracy as well. ExpO-1D-Fidelity generally had a similar effect on the interpretability metrics.

We ran experiments on the ‘MSD’ dataset³ to understand the scalability of ExpO to larger tasks. On this dataset, evaluating the interpretability metrics using MAPLE was fairly slow, and hence we only evaluate them using LIME on the first 1000 testing points. Both ExpO-Fidelity and ExpO-

³The task is to predict release year of song from a set of acoustic features, treated as a regression problem as in Bloniarz et al. (2016)

1D-Fidelity improved LIME’s interpretability metrics by at least 50% and both improved the model accuracy.

A qualitative example on the UCI ‘housing’ dataset. After sampling a random point x , we use LIME to generate a local explanation at this point for a model trained without regularization (“unregularized explanation”) and for a model trained with ExpO-1D-Fidelity (“regularized explanation”). The unregularized and regularized explanations are shown in Table 2.

Quantitatively, using ExpO-1D-Fidelity decreased the LIME-NF metric from 1.15 to 0.02; *i.e.* ExpO produced a model that is much more accurately modeled by the explanation around x . Note that the regularized explanation has fewer significant coefficient (those with absolute value greater than 0.1), and hence it is simpler as the effect is attributed to fewer features. More examples, that show similar patterns, are available in the Appendix A.5.

Medical classification experiments. The ‘support2’ dataset⁴ is used for in-hospital mortality prediction. Be-

⁴<http://biostat.mc.vanderbilt.edu/wiki/>

Table 3: Uregularized model vs. the same model trained with ExpO-Fidelity or ExpO-1D-Fidelity on the UCI regression datasets. Results are shown across 20 trials (with the standard error in parenthesis). Statistically significant improvement ($p = 0.05$) due to Fidelity is denoted in bold and due to 1D-Fidelity is underlined.

Metric	Regularizer	autoppgs	communities	day [†] (10^{-3})	housing	music	winequality.red	MSD
MSE	None	0.14 (0.03)	0.49 (0.05)	1.000 (0.300)	0.14 (0.05)	0.72 (0.09)	0.65 (0.06)	0.583 (0.018)
	Fidelity	0.13 (0.02)	0.46 (0.03)	0.002 (0.002)	0.15 (0.05)	0.67 (0.09)	0.64 (0.06)	0.557 (0.0162)
	1D-Fidelity	0.13 (0.02)	0.55 (0.04)	5.800 (8.800)	0.15 (0.07)	0.74 (0.07)	0.66 (0.06)	<u>0.548 (0.0154)</u>
LIME-PF	None	0.040 (0.011)	0.100 (0.013)	1.200 (0.370)	0.14 (0.036)	0.110 (0.037)	0.0330 (0.0130)	0.116 (0.0181)
	Fidelity	0.011 (0.003)	0.080 (0.007)	0.041 (0.007)	0.057 (0.017)	0.066 (0.011)	0.0025 (0.0006)	0.0293 (0.00709)
	1D-Fidelity	<u>0.029 (0.007)</u>	<u>0.079 (0.026)</u>	0.980 (0.380)	<u>0.064 (0.017)</u>	<u>0.080 (0.039)</u>	<u>0.0029 (0.0011)</u>	<u>0.057 (0.0079)</u>
LIME-NF	None	0.041 (0.012)	0.110 (0.012)	1.20 (0.36)	0.140 (0.037)	0.112 (0.037)	0.0330 (0.0140)	0.117 (0.0178)
	Fidelity	0.011 (0.003)	0.079 (0.007)	0.04 (0.07)	0.057 (0.018)	0.066 (0.011)	0.0025 (0.0006)	0.029 (0.007)
	1D-Fidelity	<u>0.029 (0.007)</u>	<u>0.080 (0.027)</u>	1.00 (0.39)	<u>0.064 (0.017)</u>	<u>0.080 (0.039)</u>	<u>0.0029 (0.0011)</u>	<u>0.0575 (0.0079)</u>
LIME-S	None	0.0011 (0.0006)	0.022 (0.003)	0.150 (0.021)	0.0047 (0.0012)	0.0110 (0.0046)	0.00130 (0.00057)	0.0368 (0.00759)
	Fidelity	0.0001 (0.0003)	0.005 (0.001)	0.004 (0.004)	0.0012 (0.0002)	0.0023 (0.0004)	0.00007 (0.00002)	0.00171 (0.00034)
	1D-Fidelity	<u>0.0008 (0.0003)</u>	<u>0.018 (0.008)</u>	<u>0.100 (0.047)</u>	<u>0.0025 (0.0007)</u>	0.0084 (0.0052)	<u>0.00016 (0.00005)</u>	<u>0.0125 (0.00291)</u>
MAPLE-PF	None	0.0160 (0.0088)	0.16 (0.02)	1.0000 (0.3000)	0.057 (0.024)	0.17 (0.06)	0.0130 (0.0078)	—
	Fidelity	0.0014 (0.0006)	0.13 (0.01)	0.0002 (0.0003)	0.028 (0.013)	0.14 (0.03)	0.0027 (0.0010)	—
	1D-Fidelity	<u>0.0076 (0.0038)</u>	<u>0.092 (0.03)</u>	<u>0.7600 (0.3000)</u>	<u>0.027 (0.012)</u>	<u>0.13 (0.05)</u>	<u>0.0016 (0.0007)</u>	—
MAPLE-NF	None	0.0180 (0.0097)	0.31 (0.04)	1.2000 (0.3200)	0.066 (0.024)	0.18 (0.07)	0.0130 (0.0079)	—
	Fidelity	0.0015 (0.0006)	0.24 (0.05)	0.0003 (0.0004)	0.033 (0.014)	0.14 (0.03)	0.0028 (0.0010)	—
	1D-Fidelity	<u>0.0084 (0.0040)</u>	<u>0.16 (0.05)</u>	<u>0.9400 (0.3600)</u>	<u>0.032 (0.013)</u>	<u>0.14 (0.06)</u>	<u>0.0017 (0.0008)</u>	—
MAPLE-S	None	0.0150 (0.0099)	1.2 (0.2)	0.0003 (0.0008)	0.18 (0.14)	0.08 (0.06)	0.0043 (0.0020)	—
	Fidelity	0.0017 (0.0005)	0.8 (0.4)	0.0004 (0.0004)	0.10 (0.08)	0.05 (0.02)	0.0009 (0.0004)	—
	1D-Fidelity	<u>0.0077 (0.0051)</u>	<u>0.6 (0.2)</u>	1.2000 (0.6600)	<u>0.09 (0.06)</u>	<u>0.04 (0.02)</u>	<u>0.0004 (0.0002)</u>	—

[†] The relationship between inputs and targets on the ‘day’ dataset is very close to linear and hence all errors are orders of magnitude smaller than across other datasets.

cause the output layer of our models is the softmax over logits for two classes, we run each explanation system on each of the individual logits. Table 8 in the Appendix A.6 presents the results. We observe that ExpO-Fidelity improved the interpretability metrics by 50% or more. ExpO-1D-Fidelity slightly decreased accuracy and improved the interpretability metrics by at least 25%.

4.2. Stability Regularization

For this experiment, we compared a normally trained convolutional neural network on MNIST to one trained using ExpO-Stability. Then, we evaluated the quality of saliency map explanations for these models. Both N_x and N_x^{reg} where defined as $\text{Unif}(x - 0.05, x + 0.05)$. Both the normally trained model and model trained with ExpO-Stability achieved the same accuracy of 99%. This demonstrates one of the practical differences between SENN and ExpO: SENN places strict structural constraints on the network and subsequently lowers the testing accuracy to roughly 97%. Quantitatively, training the model with ExpO-Stability decreased the stability metric from 6.94 to 0.0008. Qualitatively, training the model with ExpO-Stability made the resulting saliency maps look much better by focusing them on the presence or absence of certain pen strokes (Figure 5).

Main/SupportDesc.

5. Human Subjects Experiment Results

In the previous section, we demonstrated that training a model with ExpO improves the fidelity and stability metrics, qualitatively changes the explanations by making them simpler, and slightly improves predictive accuracy. However, if we recall the illustration from Section 1, our end goal is to help Bob better understand how the model makes decisions about the price of his house. Although we would intuitively expect an explanation with better fidelity or stability to help with this, we directly test the hypothesis with a user study.

We designed an experiment to closely simulate Bob’s scenario: using a model’s explanations, iteratively modify an instance until the model’s prediction has changed by a certain amount. Specifically, we had users increase the prediction of a neural network predicting a house’s price (using the UCI ‘housing’ dataset) by using explanations and modifying different features.

We ran the experiment with 60 participants using Amazon Mechanical Turk, and each participant completed five rounds. In each round, we showed LIME’s explanation both for a model trained with ExpO and a normally trained model, and allowed users to increase or decrease each feature to change the model’s prediction. We randomize which condition was shown on each side, and while they both start at the same point, each condition is modified independently. We display them both at the same time to control for differences in difficulty between starting points and to allow

Table 4: The results of the user study. Participants took significantly less steps to complete the task using the ExpO regularized model, and thought that the ExpO regularized model’s explanations were both more useful for completing the task and better matched their expectations of how the model would change.

Condition	Steps	Usefulness	Expectation
ExpO	8.00	28	26
None	11.45	11	11
Neither	N.A.	15	17



Figure 6: An example of the interface users were given. In the shown image, the user has taken one step for Condition A and no steps for Condition B. While the user selected ‘+’ for Item 7 in Condition A, due to the condition’s low fidelity the change had the opposite effect and decreased the price. The sign of the explanation coefficient and the underlying feature information was hidden from the user in order to simplify the task and remove confounders. Instead, they were shown each feature’s estimated impact on the model’s prediction and allowed to change each feature value to either increase or decrease the model’s prediction.

users to directly compare the two conditions. See Figure 6 for further explanation of the interface.

We recorded how many feature changes, or steps, it took the participants to reach the target price range, and then followed up with a couple questions to directly understand their preferences. We asked two qualitative questions - which condition they found more useful for completing the task, and which condition better matched their expectation of how the price would change.

While the usual number of steps taken per round was between 5 and 20, we did have some extreme outliers that took up to 400 steps. We believe these were either bots or random clicking, and so we filtered out the participants that took more steps than 99% of the results, leaving us with a total of 54 of the original 60 participants.

We found that the ExpO regularized model provided both

quantitatively and qualitatively better explanations. Quantitatively, we found that participants took 8.00 steps on average with the ExpO model compared to 11.45 steps for the normal model ($p = 0.001$). As a qualitative measure, we asked users to report which explanation they preferred and why. The participants found the ExpO condition to be both more useful for completing the task ($p = 0.012$) and better aligned with their expectation of how the model would change ($p = 0.042$). The raw results can be seen in Table 4.

We additionally collected free response answers asking why participants chose a condition, which provide us with some insight into their reasoning. Most participants who selected ExpO focused on how well the explanation’s given change matched the actual prediction’s change. For example, one participant preferred ExpO because “It seemed to do what I expected more often”, while another noted that “In Condition A [non-ExpO] the predictions seemed completely unrelated to how the price actually changed”. While some participants who preferred the unregularized model cited similar reasons, many focused on how quickly they could reach the goal rather than the quality of the explanation. For example, one participant plainly stated they preferred the unregularized model because “The higher the value the easier to hit [the] goal”; another participant similarly explained that “It made the task easier to achieve”. These participants may have benefited from the random nature of the unregularized model, which can jump to the target range at unexpected times.

6. Conclusion

In this work, we explored the idea of directly regularizing black-box models to be more interpretable with respect to the fidelity and stability metrics for local explanations. We contrasted our regularizers to classical approaches for function approximation and smoothing and provided a generalization bound for them. Next, we demonstrated that our regularizers slightly improve model accuracy and improve the interpretability metrics by somewhere from 25% to orders of magnitude across a variety of problem settings and explainers. Finally, we ran an user study demonstrating that an improvement in these metrics has practical implications for the usability of the model’s explanations. We believe that potential future work may focus on three areas: (1) exploring alternative neighborhood functions, N_x^{reg} , that match those used by other black-box explanation systems, (2) exploring how to regularize for non-local interpretability metrics, and finally (3) exploring the interaction between the regularizers and the optimization process, *e.g.*, progressively changing the importance of the regularization during training or using the regularization as an additional training step.

References

- Adebayo, J., Gilmer, J., Muelly, M., Goodfellow, I., Hardt, M., and Kim, B. Sanity checks for saliency maps. In *Advances in Neural Information Processing Systems*, pp. 9525–9536, 2018.
- Al-Shedivat, M., Dubey, A., and Xing, E. P. Contextual explanation networks. *arXiv preprint arXiv:1705.10301*, 2017.
- Alvarez-Melis, D. and Jaakkola, T. Towards robust interpretability with self-explaining neural networks. In *Advances in Neural Information Processing Systems*, pp. 7785–7794, 2018a.
- Alvarez-Melis, D. and Jaakkola, T. S. On the robustness of interpretability methods. *arXiv preprint arXiv:1806.08049*, 2018b.
- Bloniarz, A., Talwalkar, A., Yu, B., and Wu, C. Supervised neighborhoods for distributed nonparametric regression. In *Artificial Intelligence and Statistics*, pp. 1450–1459, 2016.
- Caruana, R. et al. Intelligible models for healthcare: Predicting pneumonia risk and hospital 30-day readmission. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1721–1730. ACM, 2015.
- Dheeru, D. and Karra Taniskidou, E. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>.
- Ghorbani, A., Abid, A., and Zou, J. Interpretation of neural networks is fragile. *arXiv preprint arXiv:1710.10547*, 2017.
- Kim, B., Wattenberg, M., Gilmer, J., Cai, C., Wexler, J., Viegas, F., et al. Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (tcav). In *International Conference on Machine Learning*, pp. 2673–2682, 2018.
- LeCun, Y. The mnist database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>, 1998.
- Lee, G.-H., Alvarez-Melis, D., and Jaakkola, T. S. Game-theoretic interpretability for temporal modeling. *arXiv preprint arXiv:1807.00130*, 2018.
- Lee, G.-H., Jin, W., Alvarez-Melis, D., and Jaakkola, T. Functional transparency for structured data: a game-theoretic approach. In *International Conference on Machine Learning*, pp. 3723–3733, 2019.
- Lei, T., Barzilay, R., and Jaakkola, T. Rationalizing neural predictions. *arXiv preprint arXiv:1606.04155*, 2016.
- Lipton, Z. C. The mythos of model interpretability. *arXiv preprint arXiv:1606.03490*, 2016.
- Lundberg, S. M. and Lee, S.-I. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, pp. 4765–4774, 2017.
- Plumb, G., Molitor, D., and Talwalkar, A. S. Model agnostic supervised local explanations. In *Advances in Neural Information Processing Systems*, pp. 2516–2525, 2018.
- Ribeiro, M. T., Singh, S., and Guestrin, C. Why should i trust you?: Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*, pp. 1135–1144. ACM, 2016.
- Ribeiro, M. T., Singh, S., and Guestrin, C. Anchors: High-precision model-agnostic explanations. AAAI, 2018.
- Ross, A. S., Hughes, M. C., and Doshi-Velez, F. Right for the right reasons: Training differentiable models by constraining their explanations. *arXiv preprint arXiv:1703.03717*, 2017.
- Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., and Batra, D. Grad-cam: Visual explanations from deep networks via gradient-based localization. In *Proceedings of the IEEE International Conference on Computer Vision*, pp. 618–626, 2017.
- Shrikumar, A., Greenside, P., Shcherbina, A., and Kundaje, A. Not just a black box: Learning important features through propagating activation differences. *arXiv preprint arXiv:1605.01713*, 2016.
- Simonyan, K., Vedaldi, A., and Zisserman, A. Deep inside convolutional networks: Visualising image classification models and saliency maps. *arXiv preprint arXiv:1312.6034*, 2013.
- Smilkov, D., Thorat, N., Kim, B., Viégas, F., and Wattenberg, M. Smoothgrad: removing noise by adding noise. *arXiv preprint arXiv:1706.03825*, 2017.
- Sundararajan, M., Taly, A., and Yan, Q. Axiomatic attribution for deep networks. *arXiv preprint arXiv:1703.01365*, 2017.
- Wang, F. and Rudin, C. Falling rule lists. In *Artificial Intelligence and Statistics*, pp. 1013–1022, 2015.
- Zeiler, M. D. and Fergus, R. Visualizing and understanding convolutional networks. In *European conference on computer vision*, pp. 818–833. Springer, 2014.
- Zheng, S., Song, Y., Leung, T., and Goodfellow, I. Improving the robustness of deep neural networks via stability training. In *Proceedings of the IEEE conference on*

computer vision and pattern recognition, pp. 4480–4488,
2016.

A. Appendix

A.1. Human Subjects Experiment: How is the “agent” defined?

For this task (see Section 5 for details), a greedy agent can get stuck in loops. As a result, we consider an agent that is randomized with the degree of randomization being controlled by a scaling factor, λ .

Let y denote the model’s current prediction, t denote the target value, and c_i denote the explanation coefficient of feature i . Then the score of feature i , which measures how close this feature’s estimated effect is to the target change, is: $s_i = -\lambda * (|c_i| - |y - t|)$. The agent then chooses to use feature i with probability $\frac{e^{s_i}}{\sum_j e^{s_j}}$. Looking at this distribution, we see that it is uniform (*i.e.*, does not use the explanation at all) when $\lambda = 0$ and that it approaches the greedy agent as λ approaches infinity.

In Figure 2, we run a search across the value of λ to find a rough trade-off between more frequently using the information in the explanation and avoiding loops.

A.2. Comparison of ExpO to Baseline Methods

Although there are related methods to ExpO, none are necessarily appropriate to act as a baseline for comparison:

- The regularizers defined by (Lee et al., 2018; 2019) are designed for structured data types (time-series or graphs). ExpO is designed to work with semantic features or images, so there isn’t a common experiment to run.
- SENN (Alvarez-Melis & Jaakkola, 2018a) requires that the model has a specific structure. ExpO makes no assumptions about the model’s structure.
- RTFR (Ross et al., 2017) requires that we have the domain knowledge to specify that “Feature ‘i’ should not be relevant to the prediction for point ‘x’”. ExpO does not require this or any other domain knowledge.

As a result, we consider two standard regularization techniques: l_2 and l_1 regularization. These regularizers may make the network smoother or simpler (due to sparser weights), which may make it more amenable to local explanation. The results of this experiment are in Table 5; notice that neither of these regularizers had a significant effect on the interpretability metrics compared to ExpO.

A.3. Details on the Generalization of Local Linear Explanations

Here, we provide a derivation of the bound (5) on the explanation fidelity. First, we assume that local linear explanations, β_x , are obtained by solving the ordinary least squares regression problem (as given in Algorithm 2):

$$\beta_x = [X'X'^\top]^{-1} X'f(X'), \quad (6)$$

where each column of X' denotes a sample from the neighborhood N_x and $f(X')$ is a column-vector of the corresponding function values. The *expected* fidelity of the explanation β_x can be computed analytically:

$$r(f, x) = \mathbb{E}_{N_x} [f(x')^2] - \mathbb{E}_{N_x} [f(x')x']^\top \mathbb{E}_{N_x} [[x'x'^\top]]^{-1} \mathbb{E}_{N_x} [f(x')x'] \quad (7)$$

where expectation $\mathbb{E}_{N_x} [\cdot]$ is taken with respect to x' over the neighborhood N_x . Note the equality in (7) is the expected value of the squared residual between $f(x)$ and the optimal local linear explanation, which is upper-bounded by the variance of the model in the corresponding neighborhood:

$$0 \leq r(f, x) \leq \mathbb{E}_{N_x} [f(x')^2] - \mathbb{E}_{N_x} [f(x')]^2 = \text{Var}_{N_x} [f(x')] \quad (8)$$

For instance, if $f(x)$ is L -Lipschitz and the neighborhood N_x is defined a uniform distribution within a σ -ball centered at x , then the variance of $f(x)$ within the neighborhood can be further bounded by $4L^2\sigma^2$, hence $r(f, x) \leq 4L^2\sigma^2$.

For the explanations to generalize, we would like to make sure that the gap between the average fidelity on the training set and the expected fidelity is small with high probability. More formally, the following inequality should hold:

$$\mathbb{P} \left(\mathbb{E} [r(f, x)] - \frac{1}{n} \sum_{i=1}^n r(f, x_i) > \varepsilon \right) < \delta_n(\varepsilon) \quad (9)$$

Table 5: Using l_2 or l_1 regularization has very little impact on the interpretability of the learned model.

Metric	Regularizer	autompgs	communities	day	housing	music	winequality.red
MSE	None	0.14	0.49	0.001	0.14	0.72	0.65
	L2	0.13	0.47	0.00012	0.15	0.68	0.67
	L1	0.12	0.46	1.7e-05	0.15	0.68	0.67
MAPLE-PF	None	0.016	0.16	0.001	0.057	0.17	0.013
	L2	0.015	0.17	3.2e-05	0.05	0.17	0.02
	L1	0.014	0.17	1.6e-05	0.054	0.17	0.015
MAPLE-NF	None	0.018	0.31	0.0012	0.066	0.18	0.013
	L2	0.016	0.32	4.3e-05	0.058	0.17	0.021
	L1	0.016	0.32	2.6e-05	0.065	0.18	0.016
MAPLE-Stability	None	0.015	1.2	2.6e-07	0.18	0.081	0.0043
	L2	0.011	1.3	3.2e-06	0.17	0.065	0.0058
	L1	0.013	1.2	3e-07	0.21	0.072	0.004
LIME-PF	None	0.04	0.1	0.0012	0.14	0.11	0.033
	L2	0.037	0.12	0.00014	0.12	0.099	0.047
	L1	0.035	0.12	0.00017	0.13	0.1	0.034
LIME-NF	None	0.041	0.11	0.0012	0.14	0.11	0.033
	L2	0.037	0.12	0.00015	0.12	0.099	0.047
	L1	0.036	0.12	0.00018	0.13	0.1	0.034
LIME-Stability	None	0.0011	0.022	0.00015	0.0047	0.011	0.0013
	L2	0.00097	0.032	1.7e-05	0.004	0.011	0.0021
	L1	0.0012	0.03	3e-05	0.0048	0.011	0.0016

The following is a restatement of Proposition 1 with a short proof.

Proposition 3 *Let the neighborhood sampling function N_x be characterized by some parameter σ (e.g., the effective radius of a neighborhood) and the variance of the trained model $f(x)$ across all such neighborhoods be bounded by some constant $C(\sigma) > 0$. Then, the following bound holds with at least $1 - \delta$ probability:*

$$\mathbb{E}[r(f, x)] \leq \frac{1}{n} \sum_{i=1}^n r(f, x_i) + \sqrt{\frac{C^2(\sigma) \log \frac{1}{\delta}}{2n}}$$

Proof. By assumption, the variance of the model $f(x)$ is bounded in each local neighborhood specified by N_x . Then (7) implies that each residual is bounded as $0 \leq r(f, x) \leq C(\sigma)$. Applying Hoeffding’s inequality, we get:

$$\mathbb{P}\left(\mathbb{E}[r(f, x)] - \frac{1}{n} \sum_{i=1}^n r(f, x_i) > \varepsilon\right) < \exp\left\{\frac{-2n\varepsilon^2}{C^2(\sigma)}\right\}$$

Inverting the inequality gives us the bound. ■

A.4. Choosing σ for N_x and N_x^{reg}

In Figure 7, we see that the choice of σ for N_x was not critical (the value of LIME-NF only increased slightly with σ) and that this choice of σ for N_x^{reg} produced slightly more accurate and interpretable models.

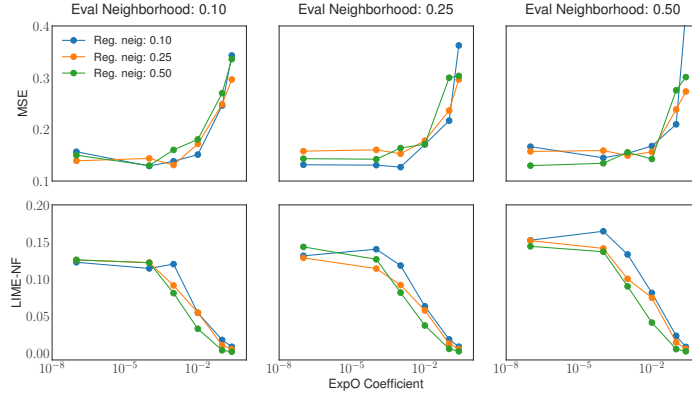


Figure 7: A comparison showing the effects of the σ parameter of N_x and N_x^{reg} on the UCI Housing dataset. The LIME-NF metric grows slowly with σ for N_x as expected. Despite being very large, using $\sigma = 0.5$ for N_x^{reg} is generally best for the LIME-NF metric and possibly for accuracy.

A.5. More Examples of ExpO’s Effects

Here, we demonstrate the effects of ExpO-Fidelity on more examples from the UCI ‘housing’ dataset (Table 6). Observe that the same general trends hold true:

- The regularized explanation more accurately reflects the model (LIME-NF metric)
- The regularized explanation generally considers fewer features to be relevant. We consider a feature to be ‘significant’ if its absolute values is 0.1 or greater.
- Neither model appears to be heavily influenced by CRIM or INDUS. The regularized model generally relies more on LSTAT and less on DIS, RAD, and TAX to make its predictions.

Table 6: More examples of how regularizing a model using ExpO-Fidelity affects the explanations. For each example we show, the feature values of the point being explained, the coefficients of the unregularized explanation, and the coefficients of the regularized explanation. Note that the bias terms have been excluded from the explanations. We also report the LIME-NF metric of each explanation.

Example Number	Value Shown	CRIM	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	LIME-NF
1	Feature	-0.36	-0.57	-0.86	-1.11	-0.14	0.95	-0.74	-1.02	-0.22	0.46	0.53	
	Unregularized Explanation	0.01	0.03	-0.14	0.31	-0.1	-0.29	0.27	-0.26	-0.07	0.13	-0.24	0.0033
	Regularized Explanation	0.0	0.01	-0.14	0.25	0.03	-0.16	0.15	-0.1	-0.12	-0.01	-0.47	0.0033
2	Feature	-0.37	-0.82	-0.82	0.66	-0.77	1.79	-0.17	-0.72	0.6	0.45	-0.42	0.0
	Unregularized Explanation	0.01	0.06	-0.15	0.32	-0.1	-0.29	0.24	-0.27	-0.12	0.11	-0.24	0.057
	Regularized Explanation	0.0	0.0	-0.15	0.25	0.01	-0.15	0.15	-0.12	-0.13	0.01	-0.47	0.00076
3	Feature	-0.35	-0.05	-0.52	-1.41	0.77	-0.13	-0.63	-0.76	0.1	0.45	1.64	
	Unregularized Explanation	-0.01	0.06	-0.16	0.29	-0.08	-0.31	0.27	-0.27	-0.11	0.1	-0.18	0.076
	Regularized Explanation	-0.03	-0.01	-0.13	0.19	-0.0	-0.15	0.14	-0.11	-0.12	0.0	-0.43	0.058
4	Feature	-0.36	-0.34	-0.26	-0.29	0.73	-0.56	-0.51	-0.12	1.14	0.44	0.14	
	Unregularized Explanation	0.02	0.06	-0.18	0.29	-0.1	-0.34	0.31	-0.21	-0.09	0.12	-0.27	0.10
	Regularized Explanation	-0.02	0.01	-0.13	0.21	0.02	-0.16	0.17	-0.11	-0.12	-0.0	-0.47	0.013
5	Feature	-0.37	-1.14	-0.88	0.45	-0.28	-0.21	-0.86	-0.76	-0.18	0.03	-0.82	
	Unregularized Explanation	0.02	0.08	-0.17	0.33	-0.11	-0.36	0.29	-0.27	-0.08	0.1	-0.28	0.099
	Regularized Explanation	-0.0	-0.0	-0.14	0.26	0.0	-0.16	0.15	-0.11	-0.15	0.01	-0.47	0.0021

The same comparison for examples from the UCI ‘winequality-red’ are in Table 7. We can see that the regularized model depends more on “volatile acidity” and less on “sulphates” while usually agreeing about the effect of “alcohol”. Further, it is better explained by those explanations than the unregularized model.

A.6. Quantitative Results on the ‘support2’ Dataset

In Table 8, we compare models trained with ExpO to those without them on the “support2” dataset.

Table 7: The same setup as Table 6, but showing examples for the UCI ‘winequality-red’ dataset

Example Number	Value Shown	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	LIME-NF
1	Feature	-0.28	1.55	-1.31	-0.02	-0.26	3.12	1.35	-0.25	0.41	-0.2	0.29	
	Unregularized Explanation	0.02	-0.11	0.14	0.08	-0.1	0.05	-0.15	-0.13	-0.01	0.31	0.29	0.021
	Regularized Explanation	0.08	-0.22	0.01	0.04	-0.04	0.06	-0.12	-0.09	-0.01	0.17	0.3	6.6e-05
2	Feature	1.86	-1.91	1.22	0.87	0.39	-1.1	-0.69	1.48	-0.22	1.96	-0.35	
	Unregularized Explanation	0.02	-0.15	0.11	0.07	-0.08	0.07	-0.23	-0.09	-0.06	0.3	0.27	0.033
	Regularized Explanation	0.09	-0.23	0.02	0.04	-0.05	0.06	-0.13	-0.09	-0.0	0.18	0.3	0.0026
3	Feature	-0.63	-0.82	0.56	0.11	-0.39	0.72	-0.11	-1.59	0.16	0.42	2.21	
	Unregularized Explanation	0.03	-0.1	0.13	0.05	-0.06	0.12	-0.21	-0.19	-0.08	0.38	0.29	0.11
	Regularized Explanation	0.09	-0.22	0.02	0.04	-0.04	0.06	-0.12	-0.09	-0.0	0.18	0.3	8.2e-05
4	Feature	-0.51	-0.66	-0.15	-0.53	-0.43	0.24	0.04	-0.56	0.35	-0.2	-0.07	
	Unregularized Explanation	0.03	-0.16	0.12	0.05	-0.13	0.09	-0.21	-0.13	-0.05	0.35	0.24	0.61
	Regularized Explanation	0.09	-0.22	0.01	0.04	-0.04	0.06	-0.12	-0.09	-0.01	0.18	0.3	6.8e-05
5	Feature	-0.28	0.43	0.1	-0.65	0.61	-0.62	-0.51	0.36	-0.35	5.6	-1.26	
	Unregularized Explanation	0.03	-0.12	0.09	0.12	-0.11	0.03	-0.19	-0.13	-0.03	0.13	0.24	0.19
	Regularized Explanation	0.08	-0.22	0.02	0.04	-0.05	0.05	-0.13	-0.09	-0.0	0.16	0.3	0.0082

Table 8: Uregularized model vs. the same model trained with ExpO-Fidelity or ExpO-1D-Fidelity on the ‘support2’ binary classification dataset. Each explanation metric was computed for both the positive and the negative class logits. Results are shown across 10 trials (with the standard error in parenthesis). Improvement due to Fidelity and 1D-Fidelity over unregularized model is statistically significant ($p = 0.05$) for all of the metrics.

Output	Regularizer	LIME-PF	LIME-NF	LIME-S	MAPLE-PF	MAPLE-NF	MAPLE-S
Positive	None	0.177 (0.063)	0.182 (0.065)	0.0255 (0.0084)	0.024 (0.008)	0.035 (0.010)	0.34 (0.06)
	Fidelity	0.050 (0.008)	0.051 (0.008)	0.0047 (0.0008)	0.013 (0.004)	0.018 (0.005)	0.13 (0.05)
	1D-Fidelity	0.082 (0.025)	0.085 (0.025)	0.0076 (0.0022)	0.019 (0.005)	0.025 (0.005)	0.16 (0.03)
Negative	None	0.198 (0.078)	0.205 (0.080)	0.0289 (0.0121)	0.028 (0.010)	0.040 (0.014)	0.37 (0.18)
	Fidelity	0.050 (0.008)	0.051 (0.008)	0.0047 (0.0008)	0.013 (0.004)	0.018 (0.005)	0.13 (0.03)
	1D-Fidelity	0.081 (0.026)	0.082 (0.027)	0.0073 (0.0021)	0.019 (0.006)	0.024 (0.007)	0.16 (0.06)

Accuracy (%): None: 83.0 ± 0.3 , Fidelity: 83.4 ± 0.4 , 1D-Fidelity: 82.0 ± 0.3 .