

Technical Report: A Stratification Approach to Partial Dependence for Codependent Variables

TERENCE PARR AND JAMES D. WILSON

University of San Francisco

September 20, 2019

Abstract

Model interpretability is important to machine learning practitioners, and a key component of interpretation is the characterization of partial dependence of the response variable on any subset of features used in the model. The two most common strategies for assessing partial dependence suffer from a number of critical weaknesses. In the first strategy, linear regression model coefficients describe how a unit change in an explanatory variable changes the response, while holding other variables constant. But, linear regression is inapplicable for high dimensional ($p > n$) data sets and is often insufficient to capture the relationship between explanatory variables and the response. In the second strategy, Partial Dependence (PD) plots and Individual Conditional Expectation (ICE) plots give biased results for the common situation of codependent variables and they rely on fitted models provided by the user. When the supplied model is a poor choice due to systematic bias or overfitting, PD/ICE plots provide little (if any) useful information.

To address these issues, we introduce a new strategy, called STRATPD, that does not depend on a user’s fitted model, provides accurate results in the presence codependent variables, and is applicable to high dimensional settings. The strategy works by stratifying a data set into groups of observations that are similar, except in the variable of interest, through the use of a decision tree. Any fluctuations of the response variable within a group is likely due to the variable of interest. We apply STRATPD to a collection of simulations and case studies to show that STRATPD is a fast, reliable, and robust method for assessing partial dependence with clear advantages over state-of-the-art methods.

Keywords: partial dependence, model interpretability, random forests, linear models, causal inference

1 Introduction

When choosing a supervised model that relates feature and response pairs, model interpretability is often at odds with predictive power. Indeed, these two objectives have traditionally led to the choice of either an interpretable *or* a predictive model (see for example Shmueli et al. (2010)). This choice has largely been divided among machine learning and statistics cultures (Breiman et al., 2001; Donoho, 2017), where machine learning practitioners focus on predictive ability and statistical practitioners focus on interpretability and inference. Recently, however, there has been a shift in the division of these two objectives as the machine learning community has begun to build what are being called “interpretable machine learning” models (Doshi-Velez and Kim, 2017; Vellido et al., 2012). Interpretable machine learning models aim to get the best of both worlds by achieving high predictive power and ensuring that the predictions of the model can be easily interpreted.

In practice machine learning model interpretation is just as important as, and in many cases more important than, obtaining a highly predictive model. Interpretable models play an essential role in artificial intelligence (Adadi and Berrada, 2018), where the interpretation of models with high predictive power like neural networks and deep learning techniques are essential to ensure robust, replicable outcomes. Interpretable models are also at the heart of applications in medicine such as precision medicine as well as psychology and psychiatry, where models are used to describe the relationship between an individual’s demographic and clinical features and their susceptibility to illness and disease, among other measureable outcomes (Dwyer et al., 2018; Katuwal and Chen, 2016).

A key component of model interpretation involves the characterization of the partial dependence of the response on any subset of the features used in the model. Consider, for example, developing a model to predict the success of a new treatment for an individual with a cognitive disorder. The partial dependence between success rate of the treatment and the age of the individual determines whether the new treatment or some alternative treatment should be used for the individual.

To describe partial dependence more formally, suppose that \mathbf{X} is an $n \times p$ matrix whose p columns represent observed features, and \mathbf{y} is $n \times 1$ vector of responses. Let X represent a randomly-selected observation (row) from \mathbf{X} and let y denote its corresponding response. Supervised algorithms seek the unknown model $f : \mathbb{R}^p \rightarrow \mathbb{R}$ that describes the relationship between X and y as $y = f(X)$. Let $C \subset \{1, \dots, p\}$ denote the index set of the features of interest and let X_C denote a randomly selected observation of these features. Let $\bar{C} = \{1, \dots, p\} \setminus C$ denote the complement of C . To assess the partial dependence of y on features X_C , one must estimate the unknown function $f_C : \mathbb{R}^p \rightarrow \mathbb{R}$ that characterizes the dependence between X_C and y :

$$y = f_C(X). \tag{1}$$

The partial dependence function f_C quantifies the dependence of y on features X_C given the data X . Estimating the partial dependence of y on any subset of variables then comes down to estimating f_C for any collection C . When X contains only one feature, a plot of the response against the feature can be used to visualize the marginal effect of the feature exactly. Given two or more features, one can similarly plot the marginal effects of each

feature separately; however, the analysis is complicated by the interactions of the variables. Generally, pairwise interaction plots are used to visualize interactions between each pair of variables; however, these are limited to pairwise analyses since interaction plots cannot directly visualize interactions along more than three dimensions (Cox and Wermuth, 2014).

To circumvent this limitation, traditional marginal plots project other axes onto the plane associated with the feature of interest and target variable. This results in marginal plots that do not isolate the specific contribution of a feature of interest to the target. For example, a marginal plot of sex (male/female) against body weight would likely show that, on average, men are heavier than women. While true, men are also taller than women on average, which likely accounts for most of the difference in average weight. It's unlikely that two "identical" people, differing only in sex, would be appreciably different in weight.

Alternatively, a linear regression model of y on the columns of \mathbf{X} provides the general trend of a single feature x_c , $c \in \{1, \dots, p\}$, on the expected value of y via the estimated regression coefficient $\hat{\beta}_c$. For a unit change in x_c , the expected value of y increases or decreases by $\hat{\beta}_c$ while holding all other variables $x_{\bar{c}}$ fixed. The major issue with fitting a linear model over the entire domain of x_c lies in the fact it does not capture non-linear relationships between (X, y) observation pairs. This is because the coefficient $\hat{\beta}_c$ is a constant, which smooths over any local y fluctuations across the entire range of x_c . Also, linear regression models are inapplicable for high dimensional ($p > n$) data sets.

Varying-coefficient models like those introduced in Fan and Zhang (2008) as well as local nonparametric methods like LOESS (Cleveland, 1979, 1981; Cleveland and Devlin, 1988) can model effect heterogeneities across the range of x_c ; however, the partial dependence between y and x_c can be difficult to interpret with these models, and these approaches are not appropriate in high dimensional settings.

The most widely used strategy to analyze partial dependence involves the combined application of partial dependence (PD) plots (Friedman, 2000) and individual conditional expectation (ICE) plots (Goldstein et al., 2015). PD and ICE both rely on the user first fitting a joint model \hat{f} for the relationship between X and y and subsequently estimate f_C through analyzing the effect that X_C has on the prediction \hat{f} . PD describes the average marginal effect of X_C , while ICE plots describe the dependence of \hat{f} on X_C . We explain the details of each of these methods in Section 4. Despite the successes of PD and ICE, there are two primary hazards of partial dependence methods that we consider in this paper:

- (i) **Model dependence:** Partial dependence depends on the fitted model \hat{f} . Thus the accuracy of such methods rely on the accuracy (and sensibility) of the chosen machine learning model. Indeed, these plots display the relationship between model prediction and features rather than the response and the features. This is problematic for several reasons. First, \hat{f} may not be a reliable model and could, for instance, sacrifice local accuracy to minimize some global loss function. Indeed, in the case that the fitted model is a poor choice, partial dependence does not provide any useful information for the user. Furthermore, the PD and ICE plots derived from different models fitted to the same (\mathbf{X}, \mathbf{y}) can look very different, making the partial dependence of y on X_C unclear.
- (ii) **Variable codependence:** Both PD and ICE, like a majority of partial dependence methods, require that the features in \mathbf{X} be pairwise independent. In practice this is

rarely the case and so the results from PD and ICE can lead to improper interpretations, as the potentially inaccurate model feeds off of potentially-nonsensical, synthesized observations arising from variable codependencies.

We explore these hazards in more detail in Sections 2 and 5. The goal of this work is to characterize partial dependence in a way that (a) does not rely on, nor make predictions from, a user’s model, and (b) does not presume mutually-independent features. We introduce a strategy, called STRATified Partial Dependence (STRATPD), that directly estimates the partial dependence of each feature without the need of a fitted model. STRATPD is a local method that first stratifies the feature space of $x_{\bar{C}}$ into disjoint regions of observations where all $x_{\bar{C}}$ variables are approximately matched across the observations in that region. The relationship between y and X_C in each $x_{\bar{C}}$ region is characterized by binning x_C itself into disjoint regions and fitting a local linear approximation through each x_C region. We describe STRATPD for numerical explanatory variables in Section 3.1 and categorical variables in Section 3.2. We apply our approach to a testbed of simulations and case studies in Section 5, finding that STRATPD provides a fast, reliable, and robust method for assessing partial dependence even in the presence of codependencies.

2 Motivation

Consider a New York City apartment rent data set from Kaggle Kaggle (2017) and the marginal plot in Figure 1(a) showing the number of bedrooms versus price. Figure 1(b) shows the (zero-centered) PD plot of rent price on the number of bedrooms as a black line. The partial dependence line is the average of the blue lines, which represent the individual conditional expectation (ICE) plots of Goldstein et al. (2015). In this case, the ICE lines depict the model prediction contributions for a single observation as the bedrooms feature shifts through all possible number of bedrooms. Because PD plots represent an average across observations, they can hide a great deal of variability, so it is helpful to combine PD and ICE plots.

While PD and ICE plots are *model-agnostic*, they are not *model-independent* and are subject to the strengths and weaknesses of the model making predictions. For example, as seen in Figure 1(b), random forests (RF) do not well extrapolate beyond their training set support range and this data set subset of 10,000 observations has very few apartments with more than 5 bedrooms. (Note the lack of blue dots in that range of the marginal plot.) PD and ICE plots shift the bedrooms feature of all observations from 0 to 8, accepting less trustworthy predictions from the model in extreme ranges.

Obtaining radically different PD and ICE plots for different underlying models is undesirable because users cannot distinguish between interesting target fluctuations and artifacts of their model choice. Consider Figure 1(c) that shows the PD/ICE plot for the exact same data set but using a Support Vector Machine (SVM with $\gamma = 1/p$). The SVM appears to have difficulty capturing the relationship between bedrooms and price, evident from the marginal plot, which means PD and ICE plots derived from an SVM for this variable are not accurate; plots derived from high-bias models should not be trusted. At the very least, users of PD and ICE should compare plots derived from multiple models.

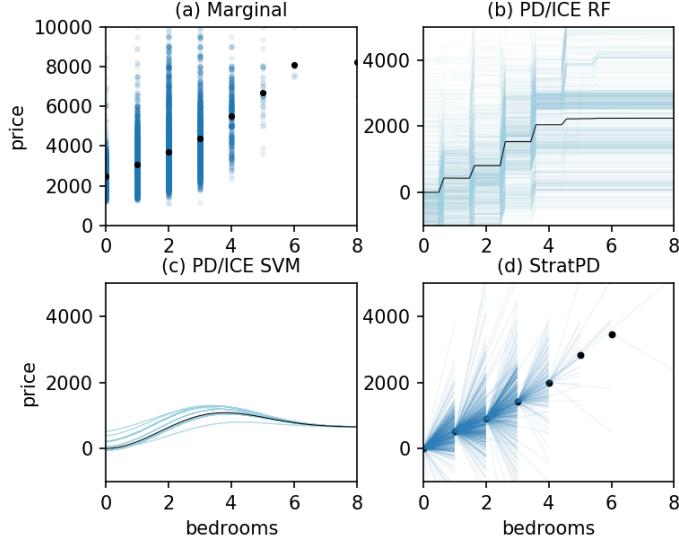


Figure 1: Plots of bedrooms versus rent price using New York City apartment rent data. (a) marginal plot, (b) PD/ICE plot derived from random forest, (c) PD/ICE plot derived from SVM, and (d) STRATPD plot; sample size is 10,000 observations of $\sim 400k$. The PD/ICE plots are radically different for the same data set, depending on the chosen user model.

Figure 1(d) shows the partial dependence of rent on the number of bedrooms (as black dots) using the STRATPD approach described in this paper. A key advantage of STRATPD is that the plot is independent of the model or models employed by the user, as STRATPD works purely from the \mathbf{X} and y relationship. The plot also depicts the density of data in the bedrooms/rent space by the number and location of lines, identifies the unique x (bedrooms) values, and characterizes the variability of the slopes across x by the spread of the line angles. The STRATPD plot does not show a data point for bedrooms=8 because there was not enough data to support any conclusions (more on this in Section 3.1). Notice that the STRATPD partial dependence plot depicts a more plausible (linear) relationship between bedrooms and rent price than either PD/ICE plot, despite the fact that the PD/ICE plots have the advantage of obtaining predictions from models, $\hat{f}(\mathbf{X})$, fitted to (\mathbf{X}, y) .

There are two remaining issues with PD plots associated with the relationship between features. First, as Friedman pointed out, PD plots are most accurate “*when [the model] is dominated by low order interactions.*” Feature interactions, such as x_1x_2 in a linear model, are difficult to tease apart to obtain partial dependence on just x_1 or x_2 . ICE plots address this issue by showing separate prediction curves for each observation as the feature of interest is moved through all possible values. This not only shows the variation hidden by the PD average curve, but it depicts interaction relationships between the feature of interest and other features.

The second issue stems from a lack of independence between features. In a nutshell, not every combination of codependent features is sensible or even possible. For example, in the apartment rent application in Figure 1, there are no apartments with five bedrooms and just one or even zero bathrooms. Similarly, there are no four bathroom studios. Because PD and ICE alter observations by shifting the feature of interest through all possible feature values,

they run the risk of conjuring up nonsensical observations that influence the calculation of partial dependence. In our experience, features in real data sets are very often codependent to some degree. This problem can be mitigated by computing PD and ICE plots on groups of mutually-dependent or interacting features of interest. For example, if sex and pregnancy are codependent features, PD/ICE can compute the partial dependence of the response variable on these two features as a single meta-feature. This would involve identifying suitable codependent feature subsets and computing PD/ICE plots for many combinations. While feasible, this approach is much harder to interpret than a single variable’s effect on the response variable.

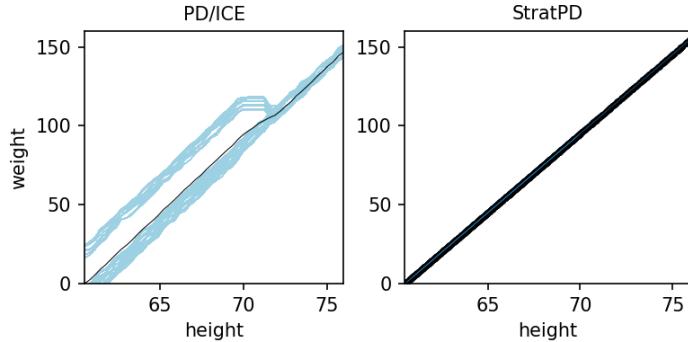


Figure 2: Plots of height versus weight using synthetic data from Equation (7). The PD/ICE on the left is biased by codependent features since pregnant women, who are typically shorter than men, have a jump in weight.

To illustrate how variable codependencies result in misleading PD and ICE plots, consider a body weight data set with observations matching a person’s characteristics to a weight in pounds. We discuss the data set details in Section 5.1 (Equation Equation (7)), but for the moment, assume that women are slightly shorter on average and are 30 pounds heavier if pregnant. The PD/ICE plot in Figure 2(a) shows an inaccurate partial dependence where shorter people appear to be slightly heavier per inch of height than those over about 72 inches. At first glance, one may surmise that there is some interesting trend between weight and individuals shorter than 72 inches. The blue ICE lines actually jump up significantly for shorter heights due to the codependence of x_{height} , x_{sex} , and $x_{pregnant}$. PD and ICE conjure up pregnant males and ask the model to estimate their weight. To be clear, the weight equation has no interaction term with x_{height} and $x_{pregnant}$, but pregnant women, who are typically shorter, have a jump in weight. The STRATPD plot in Figure 2(b), on the other hand, is not confused by codependence and gives the true partial dependence of weight on height. The next section describes the STRATPD approach and how it avoids bias from codependent variables.

3 A stratification approach

A stratification approach to estimate the partial dependence of y on x_c relies on the following two steps. First, the rows of \mathbf{X} are stratified into disjoint collections of observations for which

$x_{\bar{c}}$ are constant within each collection (ignoring x_c). Let G_j be the index set of the j th such collection of observations. The resulting pairs $\{(x_{ic}, y_i)\}_{i \in G_j}$ describe how x_c affects y , all else being equal. The second step is to fit a local linear regression model of y on x_c over each of the pairs in collection G_j :

$$y_i = \beta_{0j} + \hat{\beta}_{G_j} x_{ic} + \epsilon_i, \quad i \in G_j.$$

For each collection G_j , the estimated coefficient $\hat{\beta}_{G_j}$ quantifies the partial dependence relationship over the region $[\min(x_{ic}), \max(x_{ic})]_{i \in G_j}$. Ignoring β_{0j} (the y -intercept) removes the contribution of $x_{\bar{c}}$ to y in G_j . The regions of x_c space across collections typically overlap. In order to obtain the overall partial dependence relationship, we can partition \mathcal{X} , the domain of x_c , into disjoint regions R_1, \dots, R_m . Let $R \in \{R_1, \dots, R_m\}$ be an arbitrary region contained in \mathcal{X}) and let $\mathcal{I} = \{G_j : G_j \cap R \neq \emptyset\}$. Then, the partial dependence between x_c and y in a region R is estimated as the weighted average of all coefficients covering that region:

$$\hat{\beta}_R = \frac{1}{\sum_{G_j \in \mathcal{I}} |G_j|} \sum_{G_j \in \mathcal{I}} |G_j| \hat{\beta}_{G_j}, \quad (2)$$

where $|G|$ is the cardinality of G . The collection of regions and coefficients, $\{(R_j, \hat{\beta}_{R_j}) : j = 1, \dots, m\}$ provide a localized approximation of the partial derivative of the unknown $f(X)$ with respect to x_c .

Unfortunately, exact stratification is only feasible for two or three variables but breaks down for more variables because it is impractical to find groups of observations that are equal across so many variables. Nonetheless, stratification is simple, well understood, and clearly isolates the effect of x_c on y from the other features, even in the presence of codependent and interacting features. The only obstacle is a general and practical mechanism for stratifying observations with many variables, which leads us to the primary contribution of this paper.

3.1 StratPD for numerical variables

STRATPD seeks to isolate the local effects of x_c on the observations \mathbf{y} in the presence of confounding variables by estimating its regional effect $\hat{\beta}_R$ given in (2). Due to the challenge of exact stratification in high dimensions, STRATPD instead approximates stratification via a regression tree. The key idea is to relax stratification so that it organizes observations into groups of similar rather than equal observations. To do this, STRATPD trains a decision tree regressor, T , as in Breiman et al. (1984), on $(x_{\bar{c}}, \mathbf{y})$ to stratify observations according to the relationship between $x_{\bar{c}}$ and y . Let L_1, \dots, L_m denote the leaves of the trained tree T . Each leaf in the tree represents a region of $x_{\bar{c}}$ space and $\{(x_{ic}, y_i)\}_{i \in L_j}$ are the x_c and y observations associated with L_j . STRATPD characterizes the relationship between y and x_c on leaf L_j by following two steps:

- Partitioning $\{(x_{ic}, y_i)\}_{i \in L_j}$ into disjoint collections (bins of variable width) delimited by the unique x_c values in L_j : B_1, \dots, B_{nbins} where $nbins = |\text{unique}(x_c)| - 1$. E.g., for unique x_c values $(1, 3, 4, 5)$, the bin regions are $((1, 3), (3, 4), (4, 5))$.

⁰The STRATPD algorithm also supports the use of random forests Breiman (2001) but primarily to deal with duplicated features, as in Section 5.3.

- Fitting a simple linear regression of y on x_c from bin B_k to bin B_{k+1} to obtain $\hat{\beta}_{B_k}$.

Let $\mathcal{J} = \{B_k : B_k \cap R \neq \emptyset\}$. We estimate the partial dependence over region R using the weighted average of all $\hat{\beta}_{B_k}$ that overlap R :

$$\hat{\beta}_R^* = \frac{1}{\sum_{B_k \in \mathcal{J}} |B_k|} \sum_{B_k \in \mathcal{J}} |B_k| \hat{\beta}_{B_k} \quad (3)$$

The estimator $\hat{\beta}_R^*$ is an estimate of the partial derivative across the region R , and so numerically integrating the partial derivatives across x_c space gives a curve representing the contribution of x_c to y . Algorithm 1 encodes the STRATPD process and the full Python 3 source code is available at <https://github.com/parrt/stratx>.

Revisiting the STRATPD plot in Figure 2(d), derived from the apartment rent data set, each blue line represents a local slope $\hat{\beta}_B$ through the observations in bin B (which is a partition of x_c space whose $x_{\bar{c}}$ features are similar). Lines extend from the minimum to maximum x_c value in a specific B . Because we are interested in the relative contribution of x_c to y , STRATPD plots use zero as a y -axis baseline. The black dots represent the integration of the partial derivative estimates up to and including each unique x_c value (except the first x_c , whose integral value is 0). The partial derivative estimate at an x_c value is the (weighted) average slope of the blue lines emanating from that value. STRATPD does not interpolate between x_c points and so the plot shows dots not lines.

There are a few situations that deserve special attention. First, as the number of variables, p , increases, stratifying $x_{\bar{c}}$ into bins of similar values would normally require more and more data. But decision tree training focuses on those variables that reduce the variance in y (without using x_c), leading to leaves that best flatten y subject to hyper parameter *min_samples_leaf*. One can imagine a leaf with some variables in $x_{\bar{c}}$ that are not similar, but that would mean that training could not use them to reduce y variance within the leaf. If those variables with unequal values do not explain much y variance, they are unlikely to affect the partial dependence computation for y on x_c in L .

Second, if x_c is constant for some leaf L , then L does not support any conclusions about how changes in x_c affect y because x_c does not change in L . The observations in L , therefore, are *non-supporting observations*. Stratification leads to observations that are similar in $x_{\bar{c}}$ and, in this case, the observation x_c values are identical. That means that fluctuations in y are likely due to noise or exogenous variables not included in the data set. Such a leaf does not contribute bin coefficients, $\hat{\beta}_B$, in Algorithm 1 to the overall slope for the leaf's x_c region. This situation occurs most often when x_c contains integers.

Third, it's reasonable to ask why we use unique x_c values within each L as the bin edges, rather than splitting x_c into fixed-width bins. This was our original approach because of its simplicity, but it required another hyper parameter, such as *nbins*, and led to high proportions of non-supporting observations in some data sets. When all or most of the x_c values are discrete integers, some choices of *nbins* would lead to virtually all bins in all leaves having a single x_c value. STRATPD would generate a questionable plot due to lack of available $\hat{\beta}_B$. Also, binning integers leads to awkward bins, such as 1.8 to 3.2 bedrooms or dayofweek 1.3 to 2.1. Using the unique x_c values themselves avoids a second hyper parameter

and guarantees that bins do not have non-supporting observations unless the entire leaf has a single unique x_c .

Partial dependence through stratification also works for data sets with categorical variables in $x_{\bar{c}}$, given a suitable similarity measure between observations that supports categorical variables, but identifying an appropriate categorical similarity measure is a well-known issue. Decision trees, however, support categorical variables easily and effectively by treating categories as unique integers. Observations with categorical variables that end up in the same leaf are likely to be similar (Breiman and Cutler (2003)). Algorithm 1, therefore, works without modification for $x_{\bar{c}}$ containing categorical variables encoded as unique integers. When the column of interest, x_c , is categorical, however, a new algorithm is required.

3.2 CatStratPD for categorical variables

The stratification approach can also capture how a categorical variable x_c affects y , instead of just a single category at a time (if one were forced to dummy-encode x_c). As with STRATPD, CATSTRATPD (Algorithm 2) stratifies \mathbf{X} into groups of similar $x_{\bar{c}}$ features by training a decision tree regressor on (\mathbf{X}, \mathbf{y}) , yielding a collection of leaves. But, because categorical variables can be unordered nominal variables, the notion of y slope is not meaningful between two categories. Instead, the partial dependence plot for categorical variables shows how each category differs on average from the other categories. The category differences can be plotted as zero-centered deltas or shifted up so the lowest delta is zero.

For simplicity, we will write $(x^{(L)}, y^{(L)})$ as shorthand for the collection $\{(x_i c, y_i)\}_{i \in L}$ going forward. CATSTRATPD groups leaf observations $(x^{(L)}, y^{(L)})$ by the categories in $x^{(L)}$ and computes the average $y^{(L)}$ value per category k . To erase the y -contributions of $x_{\bar{c}}$, CATSTRATPD subtracts the average of $y^{(L)}$ from all leaf category averages, yielding a zero-centered relative increase or decrease in y for each category:

$$\begin{aligned} \mathbf{y}^{(L,k)} &= y^{(L)}[x^{(L)} = k] && (\text{Group leaf } x_c \text{ by category } k) \\ n^{(L,k)} &= |\mathbf{y}^{(L,k)}| \text{ if } |\text{unique}(x^{(L)})| > 1 \text{ else } 0 \\ \bar{y}^{(L,k)} &= \frac{1}{n^{(L,k)}} \sum_{i=1}^{n^{(L,k)}} y_i^{(L,k)} && (\text{Mean of leaf } y^{(L)} \text{ for category } k) \\ \Delta^{(L,k)} &= \bar{y}^{(L,k)} - \bar{y}^{(L)} && (\text{Remove contribution of } x_{\bar{c}} \text{ to } y^{(L)}) \end{aligned}$$

Then, to get the overall contribution of x_c to y for category k , CATSTRATPD computes the weighted average of the leaf contributions for k from all L :

$$\begin{aligned} n^{(k)} &= \sum_{T \in rf} \sum_{L \in T} n^{(L,k)} && (\text{Num supporting observations for } k) \\ \Delta^{(k)} &= \frac{1}{n^{(k)}} \sum_{T \in rf} \sum_{L \in T} n^{(L,k)} \Delta^{(L,k)} && (\text{Delta for } k \text{ is weighted average across leaves}) \end{aligned}$$

Subtracting the $x_{\bar{c}}$ contribution, $\bar{y}^{(L)}$, normalizes all $\Delta^{(L,k)}$ so they are relative to 0, providing a common baseline from which to average contributions across leaves. When there is only one category in L , the leaf does not support any conclusions about changes in y for any category so $n^{(L,k)} = 0$, dropping it from the weighted average.

Plotting category k against $\Delta^{(k)}$ gives a zero-centered partial dependence plot. The `stratx` package shifts all $\Delta^{(k)}$ so the minimum $\Delta^{(k)}$ is 0 as we find it easier to interpret the

plots. Consider the synthetic weather data shown in Figure 3(a) where temperature varies in sinusoidal fashion over the year and with different baseline temperatures per state:

$$y = t[x_{state}] + \sin\left(\frac{2\pi}{365}x_{dayofyear} + \frac{365}{2}\right) \times \epsilon, \quad \epsilon \sim N(\mu = -5, \sigma = 5) \quad (4)$$

where the baseline t per state is $\{CA = 70, CO = 40, AZ = 90, WA = 60\}$. Each sinusoid in Figure 3(a) is the average of three years' temperature data.

The categorical partial dependence plot for x_{state} is shown in Figure 3(b). CATSTRATPD stratifies $x_{\bar{c}} = \{x_{dayofyear}, x_{year}\}$ then groups these similar time buckets (leaves) by x_{state} and computes the average temperature per state in L . For each L , the average temperature in L is subtracted from the average temperature per state in L to get deltas, which are represented by blue dots in Figure 3(b). The overall temperature estimate per state is the average of those leaf averages, represented by a solid black dash. We use a strip plot to exhibit the variation and density of y values per category. The CATSTRATPD plot accurately identifies the baseline temperature per state, as does the PD/ICE plot in Figure 3(c).




Figure 3: Plots from synthetic data from Equation (4) for categorical x_c variables: (a) marginal plot of dayofyear versus temperature for four US states, (b) STRATPD plot, and (c) PD/ICE plot. Both STRATPD and PD/ICE capture the appropriate relationship, given the lack of codependence between state and other variables.

3.3 Partitioning feature space with trees and forests

Stratification of $(x_{\bar{c}}, y)$ into regions of similar observations is at the core of our approach so it's worth examining how STRATPD partitions feature space in more detail. The goal is to find groups of extremely similar $x_{\bar{c}}$ values in (\mathbf{X}, \mathbf{y}) so fluctuations in y are due solely to changes in x_c . Such groups yield pieces of the partial dependence of y on x_c . STRATPD looks for similar observations because, beyond a few variables, it's not possible to stratify observations by equal $x_{\bar{c}}$ values.

Inventing a new partitioning strategy is unnecessary because decision trees already exist that can tessellate $x_{\bar{c}}$ feature space into small hypervolumes of similar observations. If a hypervolume is tight enough, then $x_{\bar{c}}$ values are very similar and the slope of regression lines through unique x_c values of $(x^{(L)}, y^{(L)})$ is a good estimate of the partial derivative of $\frac{\partial y}{\partial x_c}$ across that Leaf. If the $x_{\bar{c}}$ volume for leaf L is too large, then $x_{\bar{c}}$ observations in L are not

similar enough to conclude that changes in y are due to x_c alone. By default, our Python implementation of STRATPD creates decision trees with at least 10 observations per leaf, but depending on y , post-training leaf size is unbounded. The x_c space within each leaf is then split into bins, dictated by the unique x_c values, in preparation for piecewise linear approximation.

Decision trees are known to overfit, which was the impetus for the invention of Random Forests(tm) (RF) Breiman (2001). The use of decision trees by default rather than RFs in STRATPD, therefore, seems an odd choice. Using RFs was our initial approach and the STRATPD algorithm and source code still support them. (The only change required to support RFs is to iterate over all leaves from all trees, rather than the leaves of a single tree, to collect local x_c slopes.) Decision trees are sufficient for partial dependence, however, because the goal is to understand the population described by the training set, not to make predictions; that is what models are for. The one exception is that multiple trees are needed to handle features in $x_{\bar{c}}$ that are identical or highly-correlated with x_c (see Section 5.3).

To reduce overfitting, RFs bootstrap and select split variables from a subset of all variables in order to create uncorrelated trees. But, that means increasing bias to some degree because each tree is trained on roughly 2/3 of the training data and without considering some of the variables. Because our goal is to group together *all* observations that are similar in *all* $x_{\bar{c}}$ variables, it is counterproductive to bootstrap and select variables from a subset. Because partial dependence is meant to explore existing (\mathbf{X}, \mathbf{y}) data rather than make predictions on future data, there's no point in sacrificing accuracy for the sake of generality.

For the data sets we examined during the preparation of this paper, moving from a decision tree to a random forest with various numbers of trees did not affect the partial dependence results; Figure 4 and Figure 5 show some typical results. The integrated partial derivative curves identified by the black dots do not change as the number of trees increases from left to right. In one simulation run for the rent data set, we did see a difference in the partial dependence dots for an extreme value of $x_{bedrooms}$ with very few y values, but it's unclear which plot is correct for this real data set. (The answer is unclear because the true partial dependence for a variable of an unknown function is unknown.)

The blue lines representing piecewise partial derivative estimates increase in number as the number of partitions (leaves) increases. Note that the variance of the partial derivative estimates is wider for RFs than for a single decision tree. This is expected because the decision tree leaves contain all observations in a feature space hypervolume and so the estimate will be less biased; RF leaves have at most 2/3 of the observations for the same hypervolume, the bootstrapping population size estimate. The education versus weight plots in Figure 4 illustrate this most clearly. The blue “cone” around the partial dependence dots widens as the number of trees increases.

Increasing the number of trees does not improve accuracy and increases the time complexity linearly in the number of trees, which is roughly what we see in practice. For example, using a single decision tree to partition a 10,000-observation rent data set sample and generate a plot takes .3s on our 4.0Ghz processor but about 8s for 30 trees (first row, far right of Figure 5).

Decision trees choose feature space hypervolumes that minimize the variance in y , but y is technically not needed to partition $x_{\bar{c}}$ into similar regions. Breiman and Cutler (2003) described how to use random forests in an unsupervised fashion by considering the original

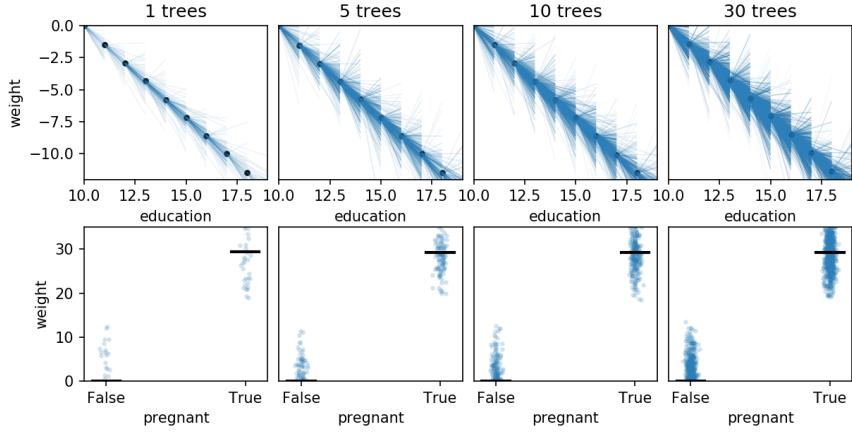


Figure 4: A comparison of single decision tree versus bootstrapped random forests with 5, 10, and 30 trees using synthetic data (2000 observations) from Equation (7). The top row shows no advantage to using more trees numerical x_c variables and the bottom row shows the same is true for categorical x_c . $\min_samples_leaf$ is 5 for education and 10 for pregnant.

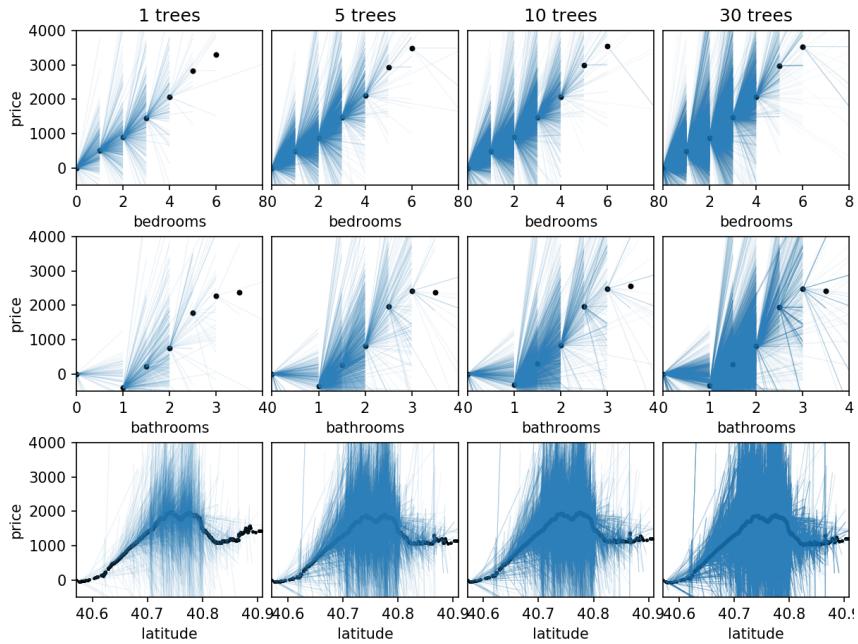


Figure 5: A comparison of single decision tree versus bootstrapped random forests with 5, 10, and 30 trees using NYC apartment rent data set. The three rows illustrate variables bedrooms, bathrooms, and latitude versus price. These graphs suggest there is no advantage to using random forests over decision trees.

\mathbf{X} matrix as class 1 and a synthesized \mathbf{X}' as class 2, which works equally well for individual decision trees, at least for this stratification application. \mathbf{X}' is just \mathbf{X} with each x'_j column bootstrapped (drawn randomly with replacement) from x_j in \mathbf{X} , effectively sampling from x_j 's marginal distribution. Figure 6 shows typical results from three variables from the rent data set and two variables from the synthesized weight data set. The left column shows unsupervised partitioning of \mathbf{X} and the right column shows the usual supervised partitioning with (\mathbf{X}, \mathbf{y}) . The results are very similar but the variance of the partial derivative estimates for the unsupervised case appears to be a bit wider. The CATSTRATPD unsupervised and supervised plots for categorical variable $x_{pregnant}$ in Figure 6(b) are virtually indistinguishable to the eye.

Figure 7 illustrates a case where unsupervised partitioning is less stable and less accurate: x_{AGE} versus x_{MEDV} (median house value) from the well-known Boston housing data set. The figure shows a marginal plot then the unsupervised and supervised STRATPD plots (and finally the PD/ICE plot). To increased stability for the unsupervised version, we used 20 trees with bootstrapping. But, in the end, there's no reason to perform unsupervised partitioning when y is always available. (Partial dependence makes no sense without y .)

The point is that partitioning $x_{\bar{C}}$ with a decision tree is more about \mathbf{X} than \mathbf{y} , which strengthens our claim of model-independence. STRATPD does not rely on a user's model, never makes predictions from internal models, and can even get away with partitioning feature space without \mathbf{y} .

4 Related Work

The PD and ICE methods mentioned in this paper each rely on the practitioner first estimating f with some machine learning model before estimating partial effects. Given a fitted model \hat{f} , PD and ICE estimate the partial effect between x_C and the fitted model \hat{f} . PD seeks to estimate the partial dependence function $\hat{f}_C^{PD}(x_C)$ given by

$$\hat{f}_C^{PD}(x_C) := \mathbb{E}_{\bar{C}} [\hat{f}(X)] = \int \hat{f}(X) d\mathbb{P}(X_{\bar{C}}).$$

The function $\hat{f}_C(x_C)$ describes the average marginal effect that the features X_C has on the prediction \hat{f} . The partial dependence of X_C is a global representation of variable dependence and averages over any heterogeneous relationships between \mathbf{y} and x_C .

The individual conditional expectation (ICE) plot from Goldstein et al. (2015) is a local method which estimates the partial dependence of the prediction $\hat{f}(X)$ on x_C across individual observations. Suppose that $(X_{C_i}, X_{\bar{C}_i})$ are the values of the i th row of \mathbf{X} . For each i , the ICE plot produces a curve from the fitted model over all values of X_C while holding $x_{\bar{C}_i}$ fixed. In particular, for observation i the following curve is plotted $\hat{f}_C^{(i)} = \hat{f}((\{X_{C_i}\}_{i=1}^n, X_{\bar{C}_i}))$. Unlike PD, ICE can be used to identify heterogeneous relationships between the fitted model and the features of interest X_C . By construction, the PD curve for a variable X_C is the average over all ICE curves for that variable. In practice, typically both ICE and PD are used to describe the partial dependence of \mathbf{y} on X_C .

An important limitation of PD and ICE is that they require independence of the features in \mathbf{X} . This is rarely the case in practice, and in such situations these plots lead to faulty

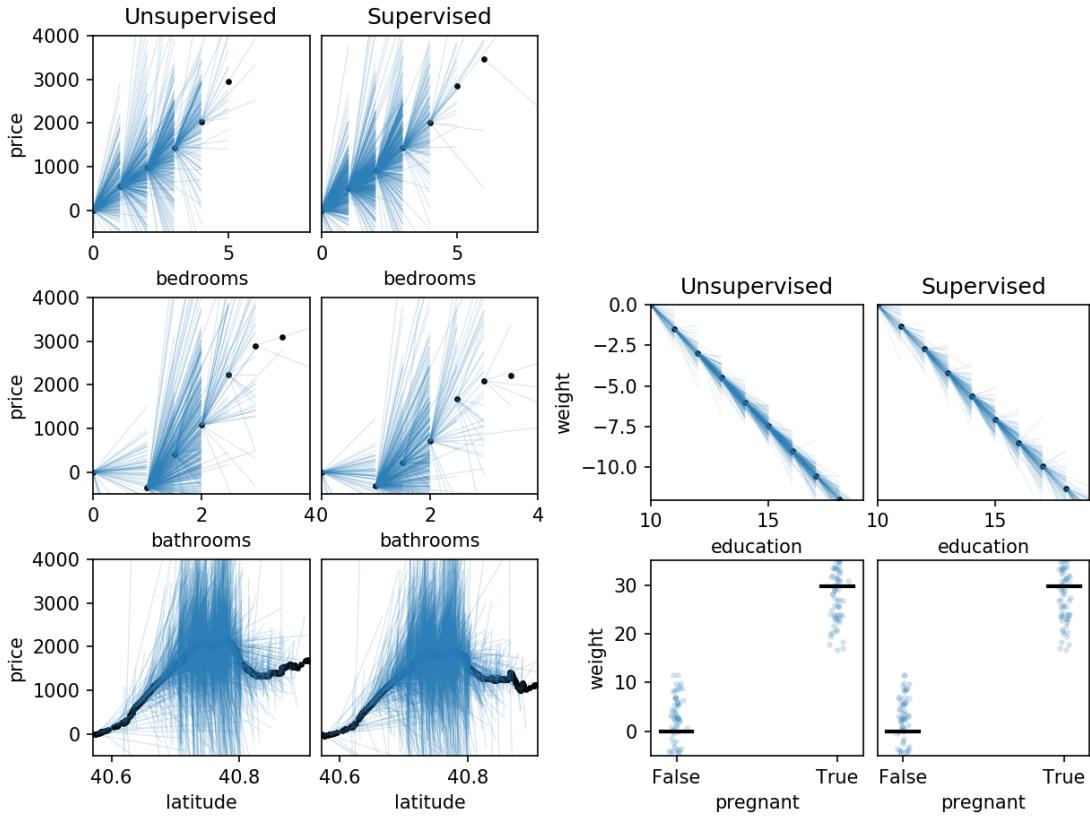


Figure 6: Comparing the effect of partitioning $x_{\bar{c}}$ space with (supervised) and without (unsupervised) y data using apartment rent data and synthetic body weight data. These graphs suggest that $x_{\bar{c}}$ can often be successfully partitioned without y .

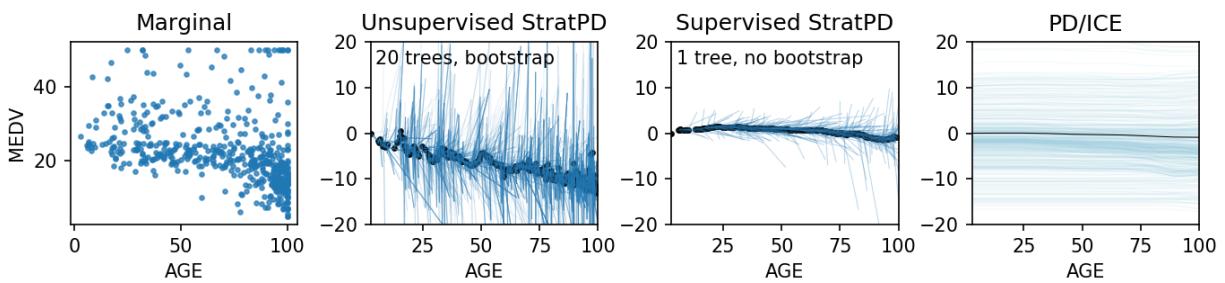


Figure 7: A demonstration that unsupervised partitioning (without y data) leads to unstable partial dependence graphs. Using a random forest instead of a single decision tree improved results, but the supervised STRATPD is still better.

inference and misinterpretation (see Apley (2016) for a discussion). Apley (2016) introduced the accumulated local effects (ALE) strategy to overcome the independence limitation. The ALE plot is an average partial dependence of X_C on \hat{f} that is calculated through the accumulation of local changes in the prediction for small windows of X_C . Local changes are measured as gradients of \hat{f} with respect to X_C while $X_{\bar{C}}$ is held fixed. Although the ALE plot is unbiased in the presence of codependent features, it still has some disadvantages. Unlike STRATPD, ALE is not directly suitable for categorical variables as an ordering of each variable is needed for the calculation of gradients of \hat{f} . Furthermore, the user must determine the number of intervals to use for calculating an ALE plot, and there is no general advice on how to do this.

Ribeiro et al. (2016) proposed the local interpretable model-agnostic explanations (LIME) method as a strategy to interpret machine learning predictions. For a prediction of interest, LIME learns an interpretable model, on a small neighborhood of data around that prediction that explains the relationship between variables and the response locally. In contrast to STRATPD, LIME is used to create local interpretable models for each prediction; however, it does not directly assess the partial dependence of the response on a subset of variables. Like LIME, STRATPD also relies upon local interpretable models; however, STRATPD does this to explain partial dependence relationships rather than correlative relationships between the response and features.

The Shapley strategy, introduced in Lundberg and Lee (2016), is a permutation-based method for estimating the relationship between y and a variable X_C through a fitted model \hat{f} . In this method, the marginal effect of X_C is represented by the Shapely value, which is the average change in the prediction made from the original data \hat{f} and the prediction made when all other variables $X_{\bar{C}}$ have been randomly shuffled in the dataset. The permuting of $X_{\bar{C}}$ is repeated many times and the average Shapely value is reported as the importance. Like PD, ICE, and ALE, the Shapley strategy is also dependent on the machine learning model fitted. Further, like any permutation method, this strategy can suffer from nonsensical observations due to the permuting of $X_{\bar{C}}$, which are subsequently incorporated in the estimated dependence. This is especially problematic in the case of highly-correlated features. Finally, a major disadvantage of the Shapley strategy is its computational complexity due to repeated permutations.

5 Experimental Results

This paper proposes a stratification approach to isolating the effect of x_c on target y and has shown a few STRATPD and CATSTRATPD plots to highlight their advantages over PD and ICE plots. In this section, we provide more examples on synthetic and real data sets, investigate the effect of noisy data, and examine how STRATPD deals with edge cases arising from unusual $x_{\bar{c}}$ partitioning. All plots, including the PD/ICE plots, were generated using the Python `stratx` library and script `genfigs.py` (in the github repository) generated all figures in this paper. (PD/ICE plots were derived from random forest models with 100 trees and minimum samples per leaf of 1).

We begin by reproducing graphs from Goldstein et al. (2015), starting with their equation in which independent variables x_2 and x_3 interact:

$$y = 0.2x_1 - 5x_2 + 10x_2 \mathbf{1}_{x_3 \geq 0} + \epsilon \quad x_1, x_2, x_3 \sim U(-1, 1), \epsilon \sim N(0, 1) \quad (5)$$

Figure 8 shows the STRATPD plots in the left column for x_2 and x_3 and the PD/ICE plots in the right column. As Goldstein et al. (2015) points out, the PD plot (top row, right side) shows no effect of x_2 on y predictions, but the ICE plot makes it clear that the apparent lack of PD effect is due to an interaction or interactions with other variables that cancel out. In this case, we know from Equation (5) that x_3 “turns off” $10x_2$ roughly half the time ($x_3 \sim U(-1, 1)$), yielding an x_2 contribution to y of $5x_2$ not $-5x_2$. The STRATPD plot also shows a relatively flat partial dependence line, although it is much smoother than the PD/ICE line. The noise, $\epsilon \sim N(0, 1)$, is about as large as the “data” and so there is nontrivial variation from simulation to simulation; curves should be taken as clues not absolute truth. Because STRATPD draws the approximate partial derivatives of y along x_2 , there is a roughly regular pattern of alternating lines of roughly slope 4 or 5, which is what we would expect since $\frac{\partial y}{\partial x_2}$ is either 5 or -5.

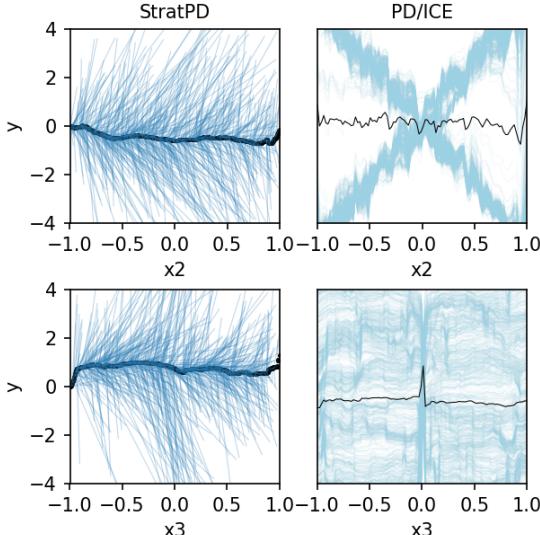


Figure 8: Plots of $x_c = x_2$ and $x_c = x_3$ from 1000 observations generated from Equation (5). STRATPD plots in the left column and PD/ICE in the right column. Both algorithms show a roughly flat partial dependence curve. PD/ICE shows a clear interaction pattern for x_2 , as does STRATPD but STRATPD’s interaction pattern is less obvious.

The partial dependence curves in the STRATPD and PD/ICE plots for x_3 in Figure 8 are also relatively flat lines because $\frac{\partial y}{\partial x_3} = 0$ when $x_3 < 0$ and $10x_2$ when $x_3 \geq 0$. Since $x_2 \sim U(-1, 1)$, $10x_2$ contributes positive and negative noise to y , which the STRATPD plot exhibits with partial derivative lines at random angles (much more random than for the

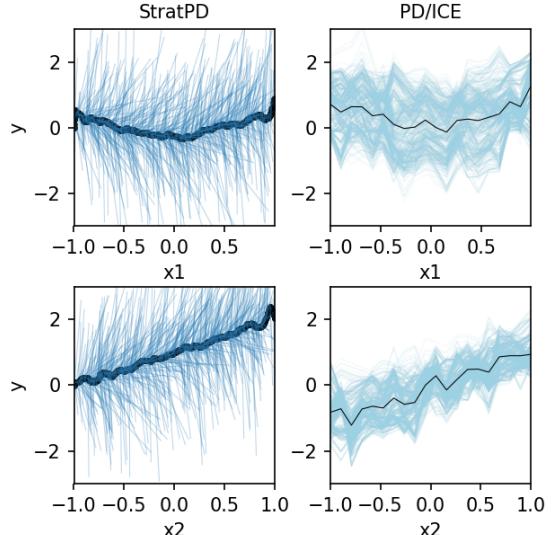


Figure 9: Plots of $x_c = x_1$ and $x_c = x_2$ from 1000 observations generated from Equation (6). STRATPD plots in the left column and PD/ICE in the right column. Both algorithms show the appropriate parabolic and linear partial dependence curves. $\text{min_samples_leaf}=10$ for both to smooth over noise.

STRATPD plot of x_2 versus y).

Goldstein et al. (2015) also demonstrate the use of ICE plots for additivity assessment using a second-order equation:

$$y = x_1^2 + x_2 + \epsilon \quad x_1, x_2 \sim U(-1, 1), \epsilon \sim N(0, 1) \quad (6)$$

We will use this equation to demonstrate that the local, nonparametric method of STRATPD can identify quadratic relationships, as shown in the left column of Figure 9; the right column shows the equivalent PD/ICE plots. The STRATPD plots for x_1 and x_2 are smoother, which can make the partial dependence relationships more pronounced visually than in the PD/ICE plots. The effect of x_2 on y should be a line with slope 1, as shown in the second row of Figure 9. Both PD/ICE and STRATPD plots give reasonable depictions of the linear relationship. To smooth out the noise, we set STRATPD hyper parameter *min_samples_leaf* to 50, which means stratification partitions $x_{\bar{c}}$ space into leaves with at least 50 observations; the default is 10. We examine noise more closely in Section 5.4.

5.1 Isolating the effect of codependent features on y

None of the variables used in Equations (5) and (6) are codependent and ICE plots have no problem exposing interactions and generating accurate PD curves. PD/ICE make the assumption that variables are independent, however, and the plots become less accurate as codependence grows. To compare STRATPD to PD/ICE for codependent variables, we synthesized a body weight data set with 2000 observations drawn from the following equation with codependence between features.

$$\begin{aligned} y &= 120 + 10(x_{height} - \min(x_{height})) + 30x_{pregnant} - 1.5x_{education} \\ \text{where } x_{sex} &\sim Bernoulli(\{M, F\}, p = 0.5) \\ x_{pregnant} &= \begin{cases} Bernoulli(\{0, 1\}, p = 0.5) & \text{if } x_{sex} = F \\ 0 & \text{if } x_{sex} = M \end{cases} \\ x_{height} &= \begin{cases} 5 * 12 + 5 + \epsilon & \text{if } x_{sex} = F, \epsilon \sim U(-4.5, 5) \\ 5 * 12 + 8 + \epsilon & \text{if } x_{sex} = M, \epsilon \sim U(-7, 8) \end{cases} \\ x_{education} &= \begin{cases} 12 + \epsilon & \text{if } x_{sex} = F, \epsilon \sim U(0, 8) \\ 10 + \epsilon & \text{if } x_{sex} = M, \epsilon \sim U(0, 8) \end{cases} \end{aligned} \quad (7)$$

Because of the codependence between $x_{pregnant}$ and x_{height} via x_{sex} , Figure 2 demonstrated that the PD/ICE plot incorrectly shows shorter people as heavier on average. PD/ICE conjures up unlikely observation such as pregnant males, resulting in biased ICE lines.

As another example of isolating codependent variable effects, compare the STRATPD and PD/ICE plots in Figure 10 showing number of years of education versus weight. Weight is related to education by slope -1.5, so a perfect partial dependence graph would show a drop of 12 pounds over 8 years of education. The PD/ICE plot captures only about two thirds of that relationship, whereas, the STRATPD plot gets the true education-weight relationship. Female observations have at least 12 years of education, versus 10 for males, so $x_{education}$ and x_{sex} are codependent, though, there is no interaction term. The fact that women are

shorter on average biases the education-weight PD/ICE plot because the baseline weight is lower from which the education contribution is subtracted.

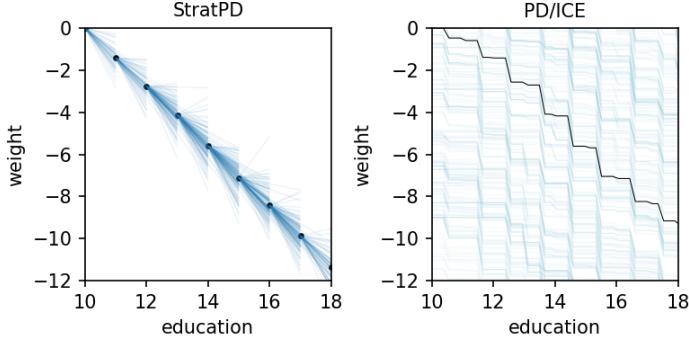


Figure 10: Plots of education versus body weight using 2000 observations from Equation (7). The STRATPD clearly identifies the linear relationship and with the proper slope of -1.5, whereas the PD/ICE plot has a more shallow slope.

Because the “signal-to-noise ratio” is low in the bodyweight data set, we set hyper parameter $\text{min_samples_leaf} = 2$ to partition $x_{\bar{c}}$ into very tight regions, leaving at least two observations from which estimate the change in y over x_c space. Tight regions increase the likelihood that y fluctuations in each leaf are due solely to changes in x_c .

5.2 The effect of model choice on PD/ICE plots

Perhaps the biggest issue with PD and ICE plots is that they rely on predictions from a user-provided model $\hat{f}(\mathbf{X})$, and different models make different assumptions and have different strengths and weaknesses. Users must choose the appropriate model for the data set and properly tune the models, otherwise ICE trendlines are untrustworthy. Figure 11 shows marginal plots, STRATPD plots, and PD/ICE plots for a 4-variable normal distribution with center $(6, 6, 6, 6)$ and covariance matrix:

$$\begin{pmatrix} 1 & 5 & .7 & 3 \\ 5 & 1 & 2 & .5 \\ .7 & 2 & 1 & 1.5 \\ 3 & .5 & 1.5 & 1 \end{pmatrix}$$

where y is related to the variables by:

$$y = x_1 + x_2 + x_3 + x_4 \quad (8)$$

The last four rows show PD/ICE plots from four different models: random forests (100 trees), support vector machines ($\gamma = 1/4$), ordinary least squares linear models, and k -nearest neighbor ($k = 5$). Because this data set is essentially skewed noise, we increased the number of data points per leaf during stratification of $x_{\bar{c}}$: $\text{min_samples_leaf} = 20$.

Because $\frac{\partial y}{\partial x_c} = 1$ for all x_c , the partial dependence curves should be lines with slope 1. The second row of Figure 11 has STRATPD plots that show linear relationships and the slopes are

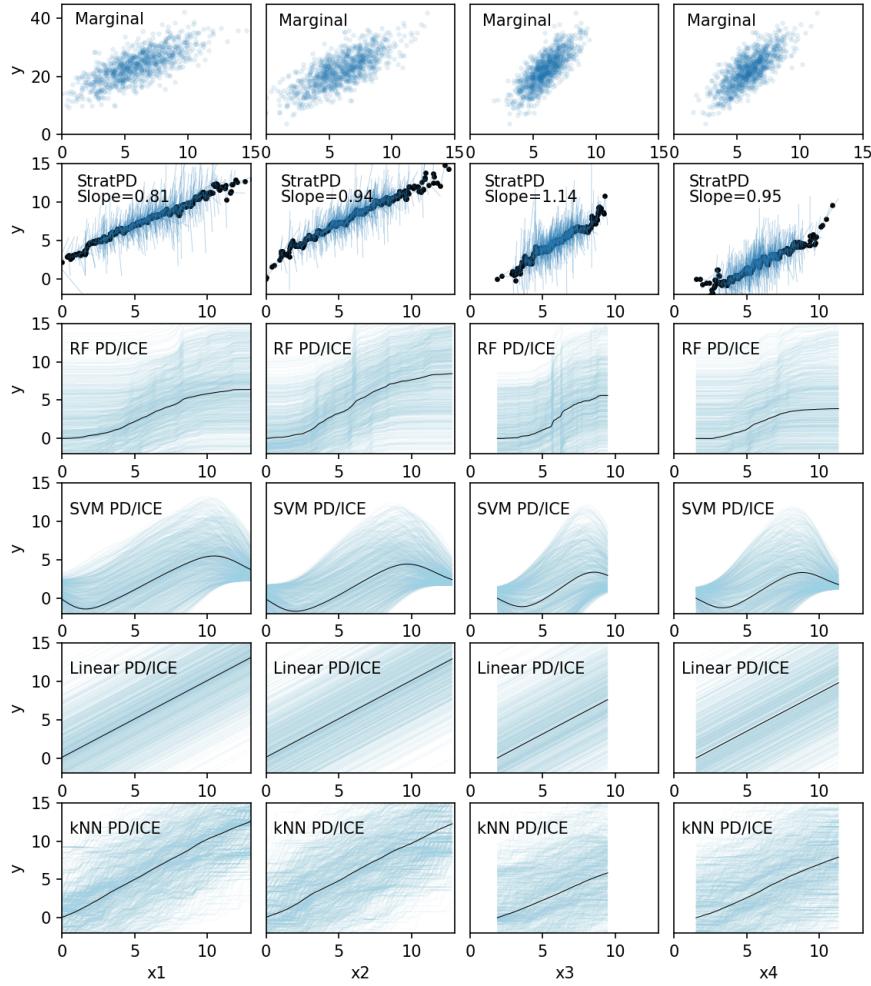


Figure 11: Marginal, STRATPD, and PD/ICE plots using 1000 operations from Equation (8) and a variety of fitted models for PD/ICE. The plots clearly demonstrate that PD/ICE results are highly dependent upon the model chosen by the user. STRATPD gets the appropriate linear partial dependence with nearly the correct slope of 1.0 for each variable.

close to 1. The PD/ICE plots derived from RF and SVM models show distinct flattening or curving behavior at the edges because the variables are codependent. Models are presented with highly unlikely combinations of x_i variables that are outside of the training data and are forced to extrapolate outside of their support range. The linear model does very well because it assumes the relationship is linear and, therefore, extrapolates linearly without issue. The nearest neighbor model also captures the linear relationship well but underestimates the partial dependence slope for x_3 and x_4 .

5.3 Duplicated columns require multiple decision trees

Isolating x_c from codependent variables in $x_{\bar{c}}$ through decision tree stratification works well in our experiments unless x_c is a linear function of a variable in $x_{\bar{c}}$. In that case, stratification hammers out variation in x_c , as if the decision tree were trained on (\mathbf{X}, \mathbf{y}) not $(x_{\bar{c}}, \mathbf{y})$. To simulate this pathological situation, we duplicated $x_{bathrooms}$ from the rent data set and generated the STRATPD and PD/ICE plots in Figure 12. The first column shows STRATPD and PD/ICE plots without a duplicated column and the second column shows the result of duplicating $x_{bathrooms}$. Both plots show much reduced partial dependence of y on the highly-predictive variable $x_{bathrooms}$. In the case of STRATPD, the stratification process groups the data by apartments with similar bathrooms and so it shows less partial dependence for the duplicated bathrooms variable. There are many fewer lines in the STRATPD plot for the duplicated column because half of the data does not support any conclusions about partial dependence: most of the leaves contain a unique x_c value.

The problem with the PD/ICE plot is not because of codependence between $x_{bathrooms}$ and its duplicate. Instead, that PD/ICE is lower because of the nature of this data set and the way random forests train. With two identical variables, decision nodes that split on one of those variables will choose between them with 50% probability. To create ICE trendlines, only one of the duplicate variables changes through $x_{bathrooms}$ space, which means that roughly half the tree decision nodes will lead to predictions that ignore the shifted x_c variable. This has the effect that the model underestimates rent prices. For example, given an observation with $x_{bathrooms} = 1$, (simplifying slightly) half the trees in the forest would predict rent appropriate for one bathroom even when the trend line shifts the x_c bathrooms to 4. It would be possible to switch models in an effort to overcome this issue, but many models do not support duplicated or highly-correlated variables.

To compensate for duplicate columns, STRATPD supports the use of random forests rather than a single decision tree to stratify $x_{\bar{c}}$ space. STRATPD uses decision trees by default rather than conventional random forests because bootstrapping is not important and, in fact, increases bias as the individual trees are working on 2/3 of the data set. The third column, top row in Figure 12 shows the STRATPD plot resulting from the use of 15 trees ($ntrees = 15$) and limiting decision node variable choice to one of two randomly-selected variables at each split ($max_split_features = 1$). Restricting the number of variables available during node splitting prevents partitioning from relying too heavily on the duplicate of x_c , leading to a number of leaves that vary in x_c .

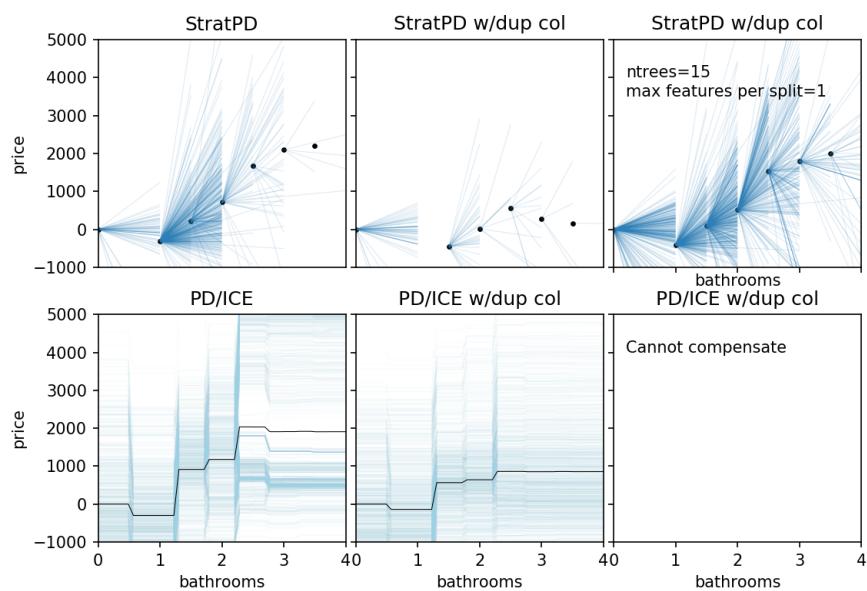


Figure 12: The effect of duplicating the predictive x_c explanatory variable using 10,000 observations of $\sim 49k$ from the rent data set. The first column shows STRATPD and PD/ICE plots without the duplicated variable and the second column shows plots after duplicating the x_c variable. To partially compensate, users can increase the number of trees but without bootstrapping and setting hyper parameter *max_split_features* to 1, as shown in the third column. PD/ICE has no way to compensate for the duplicated x_c variable.

5.4 Compensating for noise

To explore the effect of irrelevant variables on STRATPD plots, we introduced a noise column ($x_{noise} \sim U(0, 50)$) to the rent data set. Because decision trees ignore variables with low predictive power, stratification automatically ignores irrelevant or noise columns. Figure 13 shows STRATPD and PD/ICE plots for the original data set and the data set with a noise column. Both approaches are unaffected by the introduction of the noise column, but that is true for PD/ICE because this particular plot was also derived from a random forest. PD/ICE plots derived from models that are confused by noise columns would not be accurate.

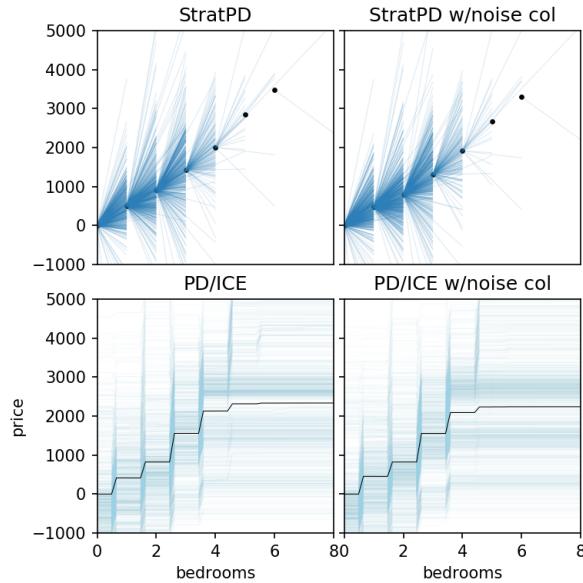


Figure 13: The effect of adding a variable of noise to \mathbf{X} using 10,000 observations of $\sim 49k$ from the rent data set. The first column shows STRATPD and PD/ICE plots without the noise variable and the second column shows plots after introducing the noise variable. The STRATPD plot ignores the noise variable as does PD/ICE because that plot was derived from a random forest model that deals well with irrelevant variables.

Overcoming noisy predictive columns or noisy y sometimes requires tuning hyper parameter $min_samples_leaf$. Figure 14 shows the effect of changing $min_samples_leaf$ on STRATPD plots at different noise levels for the Equation (6) quadratic data set. The bottom row shows decision tree partitioning of $x_{\bar{c}}=x_2$ space. The first row represents the baseline where y omits Gaussian noise. STRATPD easily picks out the quadratic relationship of depth 1 in $[-1,1]$ and is insensitive to increases in the partitioning leaf size for $x_{\bar{c}}=x_2$. As more and more Gaussian noise is added to y , the STRATPD plots become more erratic, particularly for larger partitioning leaf size as that increases the likelihood that $x_{\bar{c}}$ is influencing y . Note that the amplitude of the noise, standard deviation 1.0, in the second to last row is equal to the effect size of a parabola with depth 1.0. Partially from graphs like this, we chose the default $min_samples_leaf$ hyper parameter to be 10.

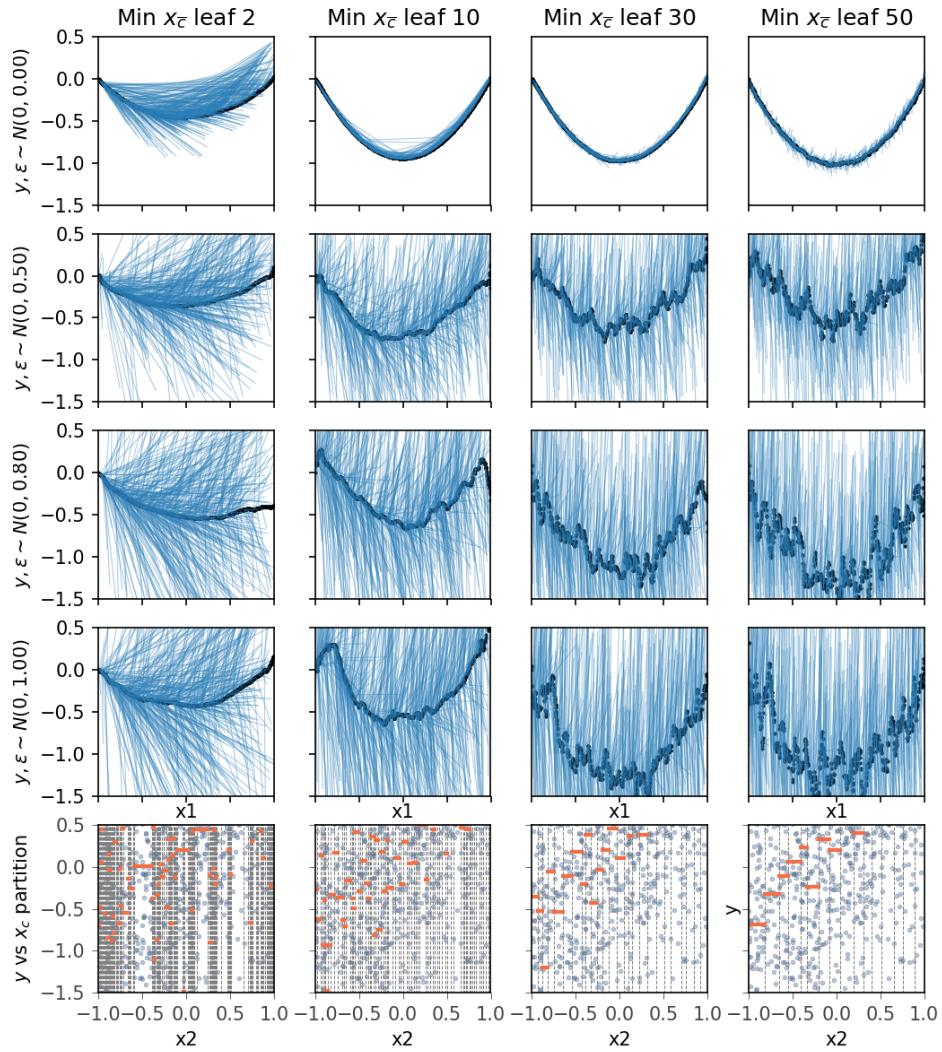


Figure 14: The effect of increasing amounts of Gaussian noise on STRATPD plots using 1000 observations from Equation (6). As the signal-to-noise ratio drops, STRATPD is less able to pick out the parabola.

5.5 StratPD and CatStratPD applied to real data

We have shown STRATPD operating on a real New York City apartment rent data set in figures such as Figure 1 and Figure 5. This section shows both STRATPD and CATSTRATPD plots for another real Kaggle data set, Kaggle (2018), concerning auction sales of used bulldozers. Of the 52 features, we selected three codependent features: YearMade, MachineHours, and ModelID. Figure 15 shows marginal plots for the three variables versus bulldozer sale price in the first column, the STRATPD and CATSTRATPD plots in the second column, and PD/ICE in the third column. The nominal variable ModelID axis in the third row was sorted by sale price. To reduce overplotting and to reduce ICE plotting time, we use the most recent 10,000 records after dropping those with missing values or zero machine hours. Random forests for PD/ICE were trained with 20 not 100 trees. The PD/ICE plot for ModelID shows just 1000 of the roughly 1500 unique values (and still takes 5 minutes to generate). the

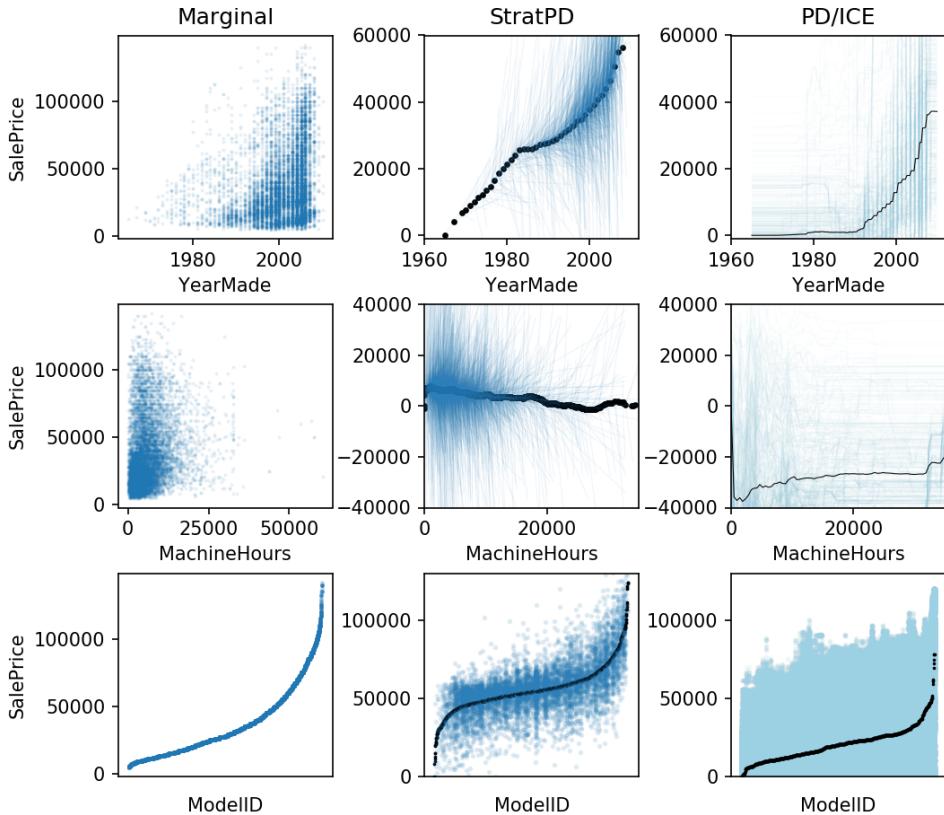


Figure 15: STRATPD and PD/ICE plots for 10,000 of ~400k observations from the bulldozer data set. The first row shows YearMade versus price, the second shows MachineHours versus price, and the third shows ModelID versus price. The STRATPD plots are plausible, certainly more plausible than the PD/ICE plots; e.g., it is unlikely that bulldozer sale prices rise as wear-and-tear (MachineHours) increases. The random forest used by PD/ICE had an out-of-bag R^2 of about 0.77 using just those three features.

As this is a real not synthesized data set, the true partial dependence curves are unknown.

Further, using only three features means the stratification approach will not be able to cancel out contributions to the sale price from the unused features. (Nontrivial feature engineering would be required to extract predictive features from the other variables and these three get an “out of bag” metric of $R^2 = 0.77$) These exogenous variables could be influencing the STRATPD and CATSTRATPD plots to be more similar to the marginal plots than is correct.

Both the STRATPD and PD/ICE plots show a price decay as bulldozers age, though the STRATPD plot is more linear for bulldozers older than about 1980 and the PD/ICE plot looks more exponential. It is possible that the linear decay shown in the STRATPD plot is more accurate because it differs significantly from the PD/ICE plot, which will suffer from the codependence of these three features. For example, the ICE lines shift the MachineHours feature into impossible observations, such as bulldozers that have been in use before they were manufactured or bulldozers sold before their ModelID existed. It is clear from both plots that there is a great deal of price variability as bulldozers age, given the STRATPD slope lines and ICE lines. The STRATPD plot (but not the PD/ICE plot) illustrates that there were many fewer bulldozers for sale that were manufactured before 1990 given the scarcity of lines in that range (which it is consistent with the marginal plot).

The STRATPD plot for MachineHours is linear with a slight negative slope in the x_c region where there is sufficient data, which is plausible. The slope lines indicate a high degree of variability that often cancel out, very much like the “big X” pattern in Figure 9. The PD/ICE plot shows a gradual increase in price as bulldozers get more use, which is highly unlikely, and also shows an immediate drop of about \$30,000 for machines that get used for a few hours. Given that the average bulldozer price is about \$36,000, that initial drop is likely due to variable codependence rather than bulldozers truly losing (then regaining) most of their value immediately.

The CATSTRATPD plot for ModelID (sorted by sale price) closely matches the marginal plot, which could be the true relationship or due to variables omitted from $x_{\bar{c}}$ or even omitted from \mathbf{X} . The PD/ICE plot also shows that some models are much more expensive than others, but is much more curvilinear.

5.6 Pathological partitioning issues

There is a pathological case to consider during $x_{\bar{c}}$ partitioning when training yields a decision tree with very large leaves, perhaps hundreds or thousands of observations. This can happen when $x_{\bar{c}}$ contains a single categorical variable or when the only strongly-predictive variable in $x_{\bar{c}}$ is categorical. The weather data set from Equation 4 is a case in point. (See the marginal plot in Figure 16(a).) Choosing $x_c = x_{dayofyear}$, means $x_{\bar{c}} = \{x_{state}, x_{year}\}$ and categorical x_{state} accounts for the largest changes in temperature. A decision tree splitting on just x_{state} , for example, would group all 365 daily temperature observations for a single state into just one leaf. (The marginal plot shows the complete sine waves but from the side, edge on.) The STRATPD plot in Figure 16(b) clearly shows the sinusoidal temperature fluctuations over the year while holding the state and year constant. The PD/ICE plot also identifies the noisy sine waves, as shown in Figure 16(c).

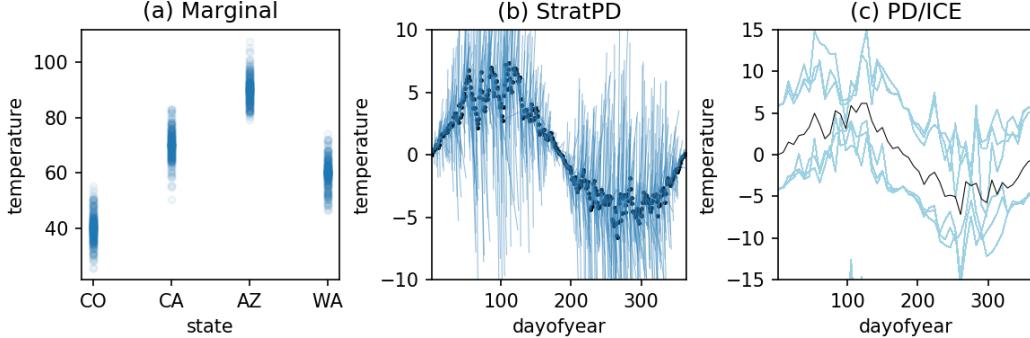


Figure 16: Pathological partitioning of x_c space can leave extremely large leaves. For x_c of dayofyear, STRATPD partitions the state variable, leaving leaves with at least a year of temperature data. STRATPD and PD/ICE plots both identify the sinusoidal relationship from three years of data using Equation (4).

5.7 Hyper parameter tuning

The key idea behind STRATPD is stratification and so hyper parameter `min_samples_leaf` is important to the operation of the algorithm; the default is 10. Larger values lead to more observations in each x_c bin, which gets more accurate $\hat{\beta}_B$ estimates. As `min_samples_leaf` gets larger, however, it is more likely that variables from $x_{\bar{c}}$ are contributing to y . Smaller values lead to much more confidence that fluctuations in y are due solely to x_c , but smaller bins lead to higher variance among the $\hat{\beta}$ estimates. Fewer observations per leaf can also cause stratification to miss nonlinear and complex behavior between y and x_c .

We recommend that users compare STRATPD plots using a number of different values of hyper `min_samples_leaf`; the `stratx` package provides function `plot_stratpd_gridsearch()` for this purpose. Figure 17 illustrates that function operating on $x_{latitude}$ versus apartment price for the rent data set. Because the partial dependence curve is fairly stable in shape and amplitude, Figure 17 increases confidence in the depicted relationship between $x_{latitude}$ and y . The “ignored” percentage shown for each graph dictates how many of the observations do not support conclusions about x_c ’s effect on y . This occurs when all x_c values within a bin are the same. As the size of the leaves grows, the size of the bins grows, which reduces the likelihood that all x_c values will be the same. The number of bins also changes the percentage of nonsupporting observations.

Figure 18 shows the results of `plot_stratpd_gridsearch()` operating on $x_{bathrooms}$ versus rent price. The partial dependence curve is stable, which again provides confidence in that partial dependence relationship.

CATSTRATPD also uses hyper parameter `min_samples_leaf` and Figure 19 shows x_{state} versus temperature for the weather data set across multiple `min_samples_leaf` values. The graphs show that when the leaf sizes too small, CATSTRATPD underestimates the effect of $state$ on temperature. Choosing a `min_samples_leaf` (2) that is less than the number of categories (4) means that CATSTRATPD cannot consider the relationship of all categories at once. After `min_samples_leaf`=20, CATSTRATPD finds the appropriate categorical partial dependence and holds study.

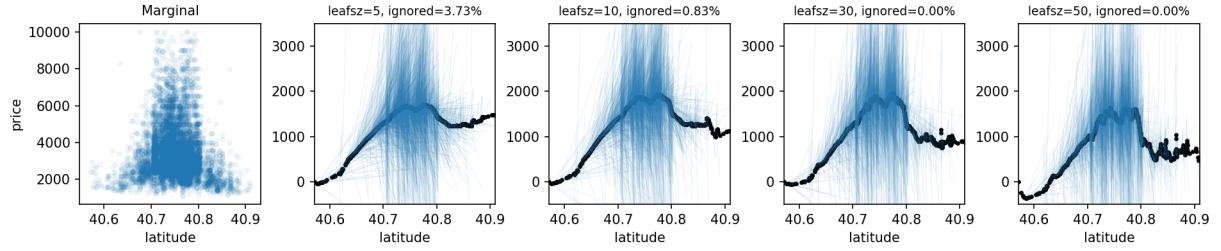


Figure 17: Hyper parameter grid for rent data showing latitude versus price; 10,000 observations.

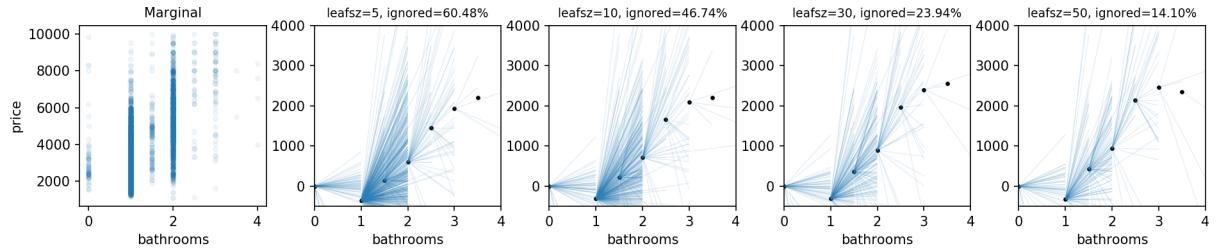


Figure 18: Hyper parameter grid for apartment rent data showing bathrooms versus price; 10,000 observations.




Figure 19: Hyper parameter grid for weather data using Equation (4) showing categorical variable state versus price. A minimum of 20 samples per leaf is required to get an accurate plot.

As discussed in Section 5.3, multiple trees are useful for dealing with duplicate columns in our experience. To engage random forest functionality, rather than simple decision tree $x_{\bar{c}}$ partitioning, there are three more hyper parameters:

1. *ntrees*. The number of trees in the random forest. The default is 1; i.e., a decision tree.
2. *max_features*. The number of features considered by the random forest when partitioning feature space. Decision trees use a default of 1.0, meaning consider all features, but random forests typically use \sqrt{p} or similar.
3. *bootstrap*. Sample $x_{\bar{c}}$ to get n observations with replacement. Decision trees use a default of false, so this must be set to true to get a random forest.

6 Discussion and Future Work

In this paper, we present an approach to partial dependence that addresses the issues encountered by other solutions. Linear regression models are not always powerful enough and cannot be used when $p > n$; PD and ICE plots are biased in the presence of codependent variables (see Figure 2, Figure 10) and yield sometimes radically different results for the same data set, depending on the model chosen by the user (see Figure 1, Figure 11).

The STRATPD approach is to stratify a data set into groups of observations that are similar in $x_{\bar{c}}$ through the use of a decision tree or random forest. Any fluctuations of y within a group (leaf) are most likely due to x_c , which lets us consider just x_c 's effect on y while holding $x_{\bar{c}}$ constant. We characterize the relationship of x_c to y within a group by fitting linear models through the unique x_c values of that leaf. To combine the resulting coefficients across the entire x_c space and across leaves, STRATPD takes the weighted average of all coefficients that overlap region R of the full x_c space; in our implementation, the R regions are taken to be the unique x_c values. The combined slope coefficients give an overall approximation to the partial derivative of y with respect to x_c and the numerical integration gives the partial dependence plot. Categorical variables are handled with a different but similar algorithm that groups y by x_c category, computes the average y for each category, and subtracts the overall average of y in that leaf to strip away contributions from $x_{\bar{c}}$. We provide algorithms for numerical and categorical data in Section 7 and Python 3 source code at <https://github.com/parrt/stratx>.

The most important advantage of STRATPD is that it directly characterizes the relationship between y and x_c *without* relying on a user's fallible model. In this way STRATPD is *model independent*, rather than *model agnostic*, meaning that it will characterize marginal relationships the same way no matter the user's choice of machine learning algorithm. While STRATPD trains a decision tree internally, the algorithm does so merely to partition feature space, never to make predictions with the tree. Surprisingly, STRATPD does not even need y to partition $x_{\bar{c}}$ space, as we demonstrated in Section 3.3 (see Figure 6), which strengthens our claim of model independence.

Another advantage over previous techniques is that STRATPD isolates the contribution of x_c to y well, at least on the data sets we've tested, even in the presence of highly-codependent variables. This is true for the synthetic data sets, where we know the answer, and for two real data sets: NYC apartment rent prices and bulldozer sales. The rent and bulldozer

plots are plausible and, moreover, are much more likely than those given by PD and ICE (see Figure 1, Figure 15). The results from PD and ICE are questionable anyway because different $\hat{f}(X)$ models on the same data sets give very different results.

Our approach has just one primary hyper parameter, *min_samples_leaf*, the minimum leaf size to use during decision tree construction. (To engage the random forest functionality, there are other hyper parameters per Section 5.6.) Our default for *min_samples_leaf* is 10 observations per leaf and the minimum *min_samples_leaf* is two in order to fit a localized linear model. Generally speaking, larger leaf sizes are desirable because they are able to capture more nonlinearities and are less susceptible to noise. As the leaf size grows, however, one risks introducing contributions from x_c into the relationship between x_c and y . We recommend examining the plots for multiple values of *min_samples_leaf*. The more consistent the plot across *min_samples_leaf* values, the higher the confidence we have in the results. For example, we provided evidence of stability in the hyper parameter with Figure 17, Figure 18, Figure 19.

Experience applying STRATPD suggests two limitations. First, STRATPD can only hold constant variables included in the \mathbf{X} matrix presented to it, but of course this is true for any partial dependence method. For example, Figure 15 is likely biased towards a traditional marginal plot than is necessary, because we included only three codependent features in \mathbf{X} for demonstration purposes.

Second, STRATPD appears to be more sensitive to noisy variables than PD and ICE, at least for the low signal-to-noise ratios found in Equation (5) and Equation (6). PD and ICE have a distinct advantage because they make use of a fitted model, $\hat{f}(X)$. Consider that PD and ICE plots derived from linear $\hat{f}(X)$ models are restricted to lines, which gives them an advantage if the underlying relationship is linear. Of course, the true $f(X)$ relationship is unknown and so choosing the right model is critical. Even nonparametric models, such as random forests, have a predetermined fit (“shape”) for x_c versus y . For $x_c = x_1$ in Equation (6), a random forest $\hat{f}(X)$ will perform generate (albeit stairstepped) parabola-like curves as PD and ICE algorithms shift x_1 through range $[-1, 1]$. We did not observe problems with STRATPD for the real data sets, which definitely have noise, but a much higher signal-to-noise ratio than the noisy synthetic data sets. As opposed to noisy predictive variables, adding superfluous noise variables does not confuse or change the STRATPD plot. This is true for PD and ICE as well, as long as the user chooses a model, such as random forests, that ignore superfluous variables. A word of caution concerning STRATPD plots. While STRATPD identifies relationships in known synthetic data sets and gives plausible results that are stable in the hyper parameter for real data sets, small fluctuations in the plots are generally not meaningful. Look for the overall trend and shape of the curve.

The next step in our research is clearly to extend STRATPD to classifiers, as we have only addressed regressors so far. Along the lines suggested by Friedman (2000), a promising approach to partial dependence for classifiers would be to swap out STRATPD’s localized linear regression models for localized logistic regression (one-versus-rest) models.

7 Algorithms

```

1 Algorithm 1: StratPD
2 Input:  $\mathbf{X}, \mathbf{y}, c,$ 
    $ntrees = 1, bootstrap = false, max\_split\_features = all,$ 
    $min\_samples\_leaf, nbins$ 
3 Output: collection of  $\beta_R$  coefficients across  $x_c$ , partial dependence curve
4 Train random forest regressor  $rf$  on  $(\bar{x}, \mathbf{y})$  with hyper-parameters:
5    $ntrees, bootstrap, max\_split\_features, min\_samples\_leaf$ 
6 foreach tree  $T \in rf$  do
7   foreach leaf  $L \in T$  do
8      $(x^{(L)}, y^{(L)}) = \{(x_{ic}, y_i)\}_{i \in L}$ 
9      $uniqx^{(L)} = \text{sorted}(\text{unique}(x^{(L)}))$ 
10    if  $|uniqx^{(L)}| > 1$  then
11       $bins = \text{split } x^{(L)} \text{ into bins delimited by range } [uniqx}_i^{(L), uniqx}_{i+1}^{(L)]}$ 
12      foreach bin  $B \in bins$  do
13         $(x^{(B)}, y^{(B)}) = \{(x_{ic}, y_i)\}_{i \in B}$ 
14         $R^{(B)} = [\min(x^{(B)}), \max(x^{(B)})]$ 
15         $n^{(B)} = |B|$ 
16        Fit linear model to  $(B_x, B_y)$  giving  $\beta_B$ 
17      end
18    end
19  end
20  $n = \sum_{T \in rf} \sum_{L \in T} \sum_{B \in L} n^{(B)}$  (Num obs.'s supporting  $\beta_B$  computations)
21  $uniqx = \text{sorted}(\text{unique}(x_c))$ 
22 for  $i = 1$  to  $|uniqx| - 1$  do
23    $R = (uniqx_i, uniqx_{i+1})$ 
24   foreach bin  $B$  created above do
25      $\beta_R = \frac{1}{n} \sum_{B \in R} \{n^{(B)} \beta_B\}$ 
26   end
27 end
28  $pd = \text{numerically integrate } \beta_R \text{'s across } uniqx$ 
29 return collection of all  $\beta_R, pd$ 

```


References

- Adadi, A. and Berrada, M. (2018). Peeking inside the black-box: A survey on explainable artificial intelligence (xai). *IEEE Access*, 6:52138–52160.
- Apley, D. W. (2016). Visualizing the effects of predictor variables in black box supervised learning models. *arXiv preprint arXiv:1612.08468*.
- Breiman, L. (2001). Random forests. *Machine Learning*, 45(1):5–32.
- Breiman, L. and Cutler, A. (2003). Random forests website. https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm#unsup. Accessed: 2019-05-24.
- Breiman, L. et al. (2001). Statistical modeling: The two cultures (with comments and a rejoinder by the author). *Statistical science*, 16(3):199–231.
- Breiman, L., Friedman, J. H., Olshen, R. A., and Stone, C. J. (1984). *Classification and Regression Trees*. Wadsworth.
- Cleveland, W. S. (1979). Robust locally weighted regression and smoothing scatterplots. *Journal of the American statistical association*, 74(368):829–836.
- Cleveland, W. S. (1981). Lowess: A program for smoothing scatterplots by robust locally weighted regression. *American Statistician*, 35(1):54.
- Cleveland, W. S. and Devlin, S. J. (1988). Locally weighted regression: an approach to regression analysis by local fitting. *Journal of the American statistical association*, 83(403):596–610.
- Cox, D. R. and Wermuth, N. (2014). *Multivariate dependencies: Models, analysis and interpretation*. Chapman and Hall/CRC.
- Donoho, D. (2017). 50 years of data science. *Journal of Computational and Graphical Statistics*, 26(4):745–766.
- Doshi-Velez, F. and Kim, B. (2017). Towards a rigorous science of interpretable machine learning. *arXiv preprint arXiv:1702.08608*.
- Dwyer, D. B., Falkai, P., and Koutsouleris, N. (2018). Machine learning approaches for clinical psychology and psychiatry. *Annual review of clinical psychology*, 14:91–118.
- Fan, J. and Zhang, W. (2008). Statistical methods with varying coefficient models. *Statistics and its Interface*, 1(1):179.
- Friedman, J. H. (2000). Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29:1189–1232.
- Goldstein, A., Kapelner, A., Bleich, J., and Pitkin, E. (2015). Peeking inside the black box: Visualizing statistical learning with plots of individual conditional expectation. *Journal of Computational and Graphical Statistics*, 24(1):44–65.

- Kaggle (2017). Two sigma connect: Rental listing inquiries. <https://www.kaggle.com/c/two-sigma-connect-rental-listing-inquiries>. Accessed: 2019-06-15.
- Kaggle (2018). Blue book for bulldozer. <https://www.kaggle.com/sureshsubramaniam/blue-book-for-bulldozer-kaggle-competition>. Accessed: 2019-06-15.
- Katuwal, G. J. and Chen, R. (2016). Machine learning model interpretability for precision medicine. *arXiv preprint arXiv:1610.09045*.
- Lundberg, S. and Lee, S.-I. (2016). An unexpected unity among methods for interpreting model predictions. *arXiv preprint arXiv:1611.07478*.
- Ribeiro, M. T., Singh, S., and Guestrin, C. (2016). Why should i trust you?: Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*, pages 1135–1144. ACM.
- Shmueli, G. et al. (2010). To explain or to predict? *Statistical science*, 25(3):289–310.
- Vellido, A., Martín-Guerrero, J. D., and Lisboa, P. J. (2012). Making machine learning models interpretable. In *ESANN*, volume 12, pages 163–172. Citeseer.